

COL 362 & COL 632

SQL Recursion and Constraints

27 Jan 2023

relvars

With Clause

max-budget

value
max(budget)

- The **with** clause provides a way of defining a temporary relation whose definition is available only to the query in which the **with** clause occurs.
- Find all departments with the maximum budget

`department(dept_name, building, budget)`

Also called
Common Table
Expressions
(CTEs)

```
with max_budget (value) as
    (select max(budget) from department)
select department.name
from department, max_budget
where department.budget = max_budget.value;
```

max-budget

Complex Queries using With Clause

- Find all departments where the total salary is greater than the average of the total salary at all departments

`instructor(ID, name, dept_name, salary)`

```
with dept_total (dept_name, value) as
    (select dept_name, sum(salary)
     from instructor
     group by dept_name),
  dept_total_avg (value) as
    (select avg(value)
     from dept_total)
select dept_name
from dept_total, dept_total_avg
where dept_total.value > dept_total_avg.value;
```

SQL - Recursion

Fix point

`prereq(course_id, prereqCourse_id)`

find which courses are a prerequisite, whether directly or indirectly, for a specific course

```
with recursive rec_prereq(course_id, prereq_id) as (  
    select course_id, prereqCourse_id  
    from prereq  
    union  
    select rec_prereq.course_id, prereq.prereq_id  
    from rec_prereq, prereq  
    where rec_prereq.prereq_id = prereq.course_id  
)  
Select *  
from rec_prereq;
```

prereqs for a course X?

1. X's prereqs (i.e., from the *prereq* relation)
2. Courses that are prereqs of all courses that are the prereqs found in step 1

SQL - Recursion

- Recursive views make it possible to write queries, such as transitive closure queries, that cannot be written without recursion or iteration.
- **Intuition:** *Without recursion*, a non-recursive non-iterative program can perform only a fixed number of joins of *prereq* with itself
 - This can give only a fixed number of levels of prerequisites
 - Given a fixed non-recursive query, we can construct a database with a greater number of levels of prerequisites on which the query will not work
 - **Alternative:** write a procedure to iterate as many times as required

Restrictions on Recursion

- **query must be monotonic**

- Without this there is no guarantee of termination
- Termination – when you reach a **fix-point**



- SQL prohibits the use of following in recursive query definition:

- Aggregation on the recursive view.
- **not exists** on a subquery that uses the recursive view
- Set difference (**except**) whose right-hand side uses the recursive view.



- Each of these make the query size to *reduce* during the recursive computation

A Small Homework

- Design a relational schema and SQL queries to
 - i. Store the basis functional dependencies
 - ii. Check if a functional dependency is part of the closure of basis set of functional dependencies stored in your relation(s) designed above

Constraints

Constraints for DB Consistency

- Security Constraints
 - Access control mechanisms
 - Restricted views
 - Prevent unauthorized access to relations and attributes
 - *Eg. A student may not be allowed to see records of another student,*
- Integrity constraints guard against accidental damage to the database, by ensuring that authorized changes to the database do not result in a loss of data consistency
 - *A current account must have a monthly balance greater than Rs. 5,000*
 - *Minimum salary of an employee with graduate degree in Delhi should be Rs. 21,000*
 - *Every faculty member must have a (non-null) phone number*

Constraints on Single Relation

- **not null**
- **primary key**
- **unique**
- **check (P)**, where P is a predicate

Specified as part of the
relation declaration

Constraints on Single Relation

- **not null**

- Declare name and budget to be not null (as part of the relation defn.)

- name varchar(20) **not null**

- budget numeric(12,2) **not null**

- **primary key**

- **unique**

- **check (P)**, where P is a predicate

Constraints on Single Relation

- not null
- primary key
- **Unique**
unique (A1, A2, ..., Am)
 - The unique specification states that the attributes A1, A2, ..., Am form **a candidate key**
 - Candidate keys are permitted to be null (in contrast to primary keys).
- **check** (P), where P is a predicate

Constraints on Single Relation

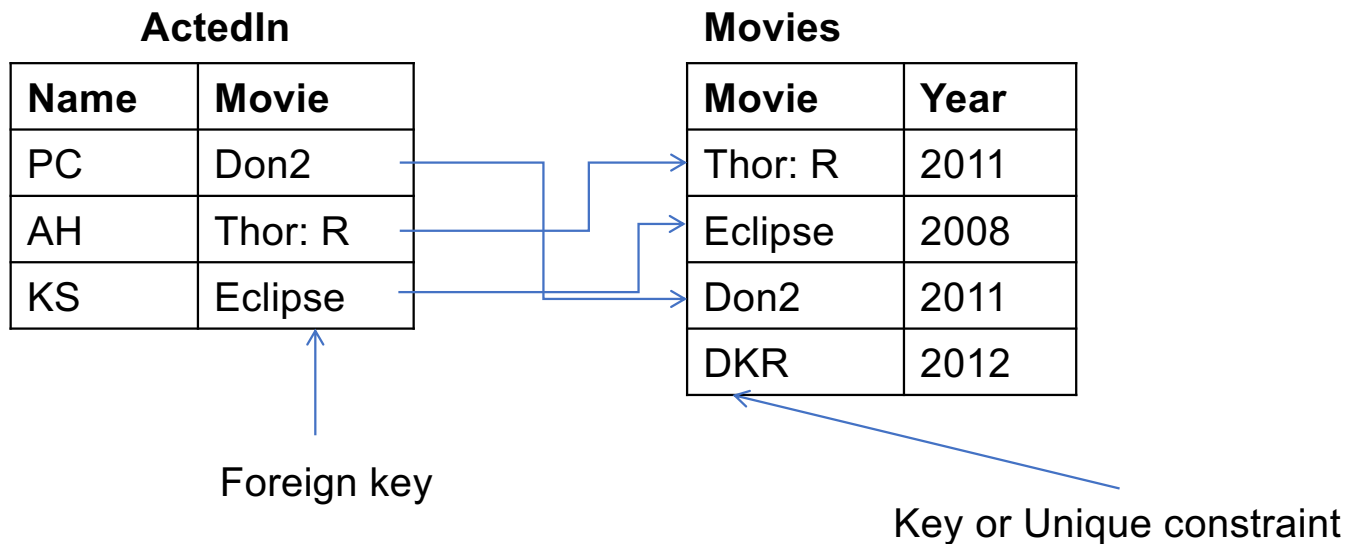
- **check (P)**, where P is a predicate

The **check (P)** clause specifies a predicate P that must be satisfied by every tuple in a relation.

Example: ensure that semester is one of fall, winter, spring or summer

```
create table section
(course_id varchar (8),
 sec_id varchar (8),
 semester varchar (6),
 year numeric (4,0),
 building varchar (15),
 room_number varchar (7),
 time slot id varchar (4),
 primary key (course_id, sec_id, semester, year),
 check (semester in ('Fall', 'Winter', 'Spring', 'Summer')))
```

Referential Integrity (1/2)



```
CREATE TABLE ActedIn (  
    Name varchar(30), Movie varchar(30),  
    CONSTRAINT actedIn_fkey  
        FOREIGN KEY (Movie) REFERENCES Movies (Movie));
```

Referential Integrity (2/2)

- Reject, Cascade, Set-null

ActedIn

Name	Movie
PC	Don2
AH	Thor: R
KS	Eclipse
CB	DKR

Movie

Movie	Year
Thor: R	2011
Eclipse	2008
Don2	2011
DKR	2012

Reject modifications which violate constraints

“Transfer” modifications

Set attribute(s) to null if needed

insert into actedIn: ('CB', 'DKR')

insert into actedIn: ('AH', 'Thor: Ragnarok')

delete from actedIn: ('KS', 'Eclipse')

delete from Movie: ('DKR', 2012)

delete from Movie: ('Eclipse', 2008)

update Movie: ('DKR', 2012) to ('DK', 2011)

```
CREATE TABLE ActedIn (  
  Name varchar(30), Movie varchar(30),  
  CONSTRAINT actedIn_fkey  
    FOREIGN KEY (Movie) REFERENCES Movie (movieid)  
    ON DELETE CASCADE)
```

“Cyclic” Constraints (1/3)

Actors

Name	Age	Addr	Famous_Movie
Priyanka Chopra	36	Mumbai	Don-II
Anthony Hopkins	81	LA	Thor: R
Bill Nighy	69	LA	Valkyrie
Abhishek Bachchan	42	Mumbai	Raavan

Movies

Name	Year	Title
Priyanka Chopra	2011	Don-II
Anthony Hopkins	2011	Thor: R
Bill Nighy	2009	Valkyrie
Abhishek Bachchan	2010	Raavan
Anthony Hopkins	2003	TLS



insert into Actors: ('Kristen Stewart', 23, 'LA', 'Breaking Dawn');
insert into Movies: ('Kristen Stewart', 2011, 'Breaking Dawn');

“Cyclic” Constraints (2/3)

- Notion of a “Transaction”

- An atomic unit of execution

The state of the database is *consistent* before and after a successful completion of a transaction

- Currently, the two inserts together form a single transaction

insert into Actors: ('Kristen Stewart', 23, 'LA', 'Breaking Dawn');

insert into Movies: ('Kristen Stewart', 2011, 'Breaking Dawn');

- **Defer** constraint checking until after transaction


“Cyclic” Constraints (3/3)

Actors

Name	Age	Addr	Famous_Movie
Priyanka Chopra	36	Mumbai	Don-II
Anthony Hopkins	81	LA	Thor: R
Bill Nighy	69	LA	Valkyrie
Abhishek Bachchan	42	Mumbai	Raavan

Movies

Name	Year	Title
Priyanka Chopra	2011	Don-II
Anthony Hopkins	2011	Thor: R
Bill Nighy	2009	Valkyrie
Abhishek Bachchan	2010	Raavan
Anthony Hopkins	2003	TLS



```
CREATE TABLE Actors (  
  Name varchar(30), Age int, Addr varchar (30), Famous_Movie varchar(30),  
  CONSTRAINT actors_fkey  
    FOREIGN KEY (Famous_Movie) REFERENCES Movies (Title)  
    DEFERRABLE INITIALLY DEFERRED)
```

Transactions

- A **transaction** consists of a sequence of query and/or update statements and is a “unit” of work
- The SQL standard specifies that a transaction begins implicitly when an SQL statement is executed.
- The transaction must end with one of the following statements:
 - **Commit work.** The updates performed by the transaction become permanent in the database.
 - **Rollback work.** All the updates performed by the SQL statements in the transaction are undone.
- Atomic transaction
 - either fully executed or rolled back as if it never occurred
- Isolation from concurrent transactions

