


# **COL 362 & COL 632**

Processing SQL – Physical Plans

14 Feb 2023

cardinality estimation

# Physical operators – Selection



$\sigma_{year > 1947}$


- **SCAN** the tuples in *Movies* one by one
- Retain tuples which satisfy the condition

- If an index is present on *Year*, then perform an **INDEX-SCAN**
- Fetch tuples which satisfy the condition

The output of the operator: disk or memory?

$[r_i, s_i]$

# Physical Operators – Joins (1/2)



When won't this work?

$\bowtie$  A.age  
 ~~$\bowtie$~~  Now  
in release

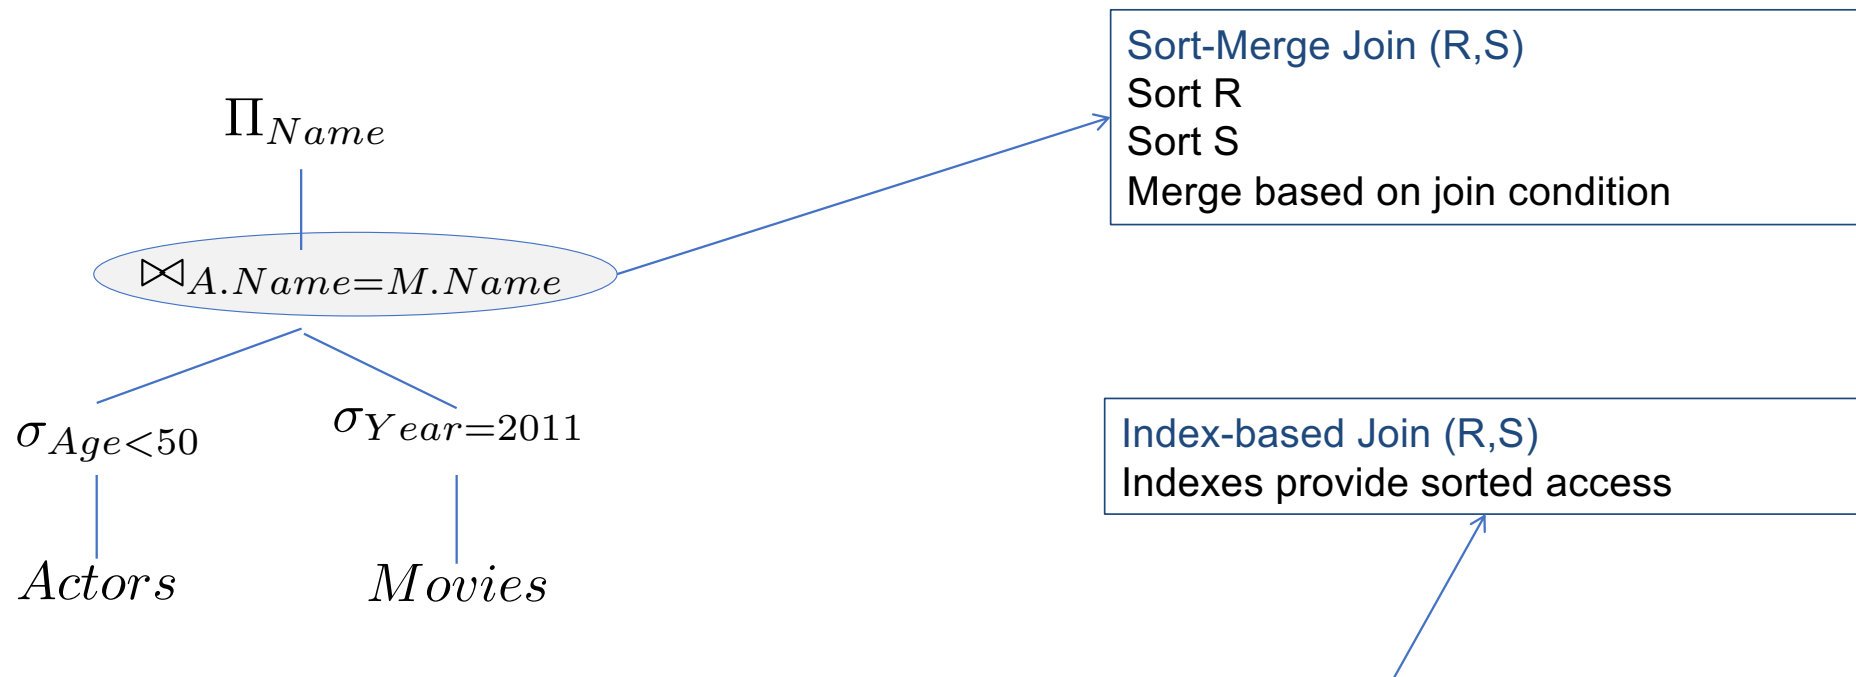
## Nested-loop Join (R,S)

```
iterate over  $r_i$  in R  
  iterate over  $s_i$  in S  
    if (join-condition satisfied)  
      output join( $r_i, s_i$ )  
    end iteration  
  end iteration  
end iteration
```

## Hash Join (R,S)


```
iterate over  $r_i$  in R  
  hash  $r_i$   
iterate over  $s_i$  in S  
  hash  $s_i$   
  join tuples in same bucket ✓
```

# Physical Operators – Joins (2/2)



What kind of indexes are ideal?

# Physical Operators – Projection



- **SCAN** the tuples one by one
- Retain required attributes

# Grouping and Aggregation

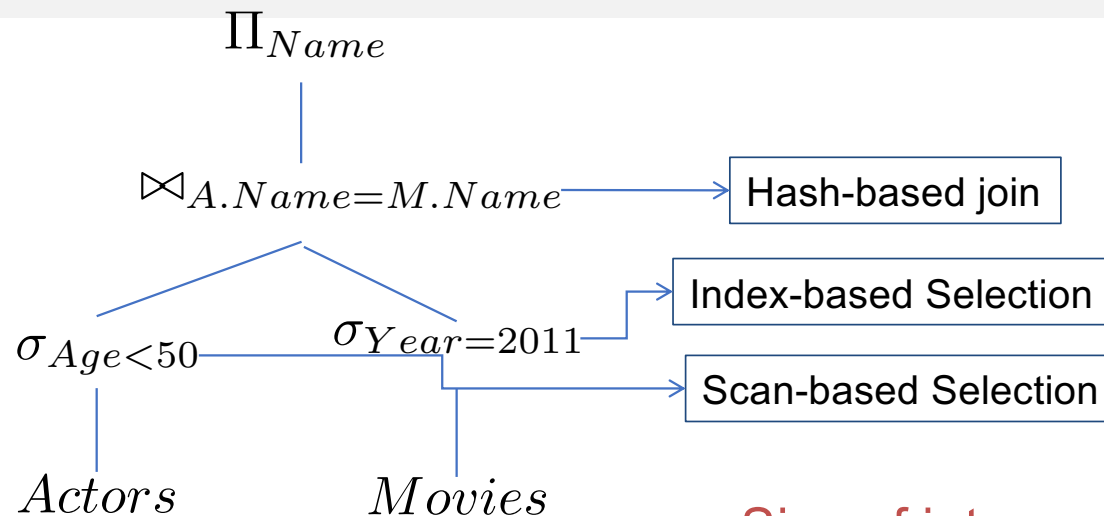
- Hash-based grouping
- Scan-based aggregation

$\gamma_{\{G,A\}}(S)$

For each group, compute aggregates A

Hash each tuple on the grouping attributes G

# Cost of a plan



- Size of intermediate results
  - Order of operations
  - Choice of algorithms
- Various system parameters
  - Available memory ✓
  - Disk contention ✓

# Cost of a Query

- Depends on where you store the data ✓
  - Memory, Disk, SSD, Tapes, Cloud, ...
- How you access the data ✓
  - Sequential vs. Random
  - Blocked transfer vs. "just what you need"
- If one has access path (or index) and its cost ✓
- CPU cost per tuple ✓
- CPU cost for each index entry (+I/O) ✓
- CPU cost per operator ✓

① Number of disk blocks read/written

② Number of "random" I/O.



# Analyzing Query Plans

- EXPLAIN statement

HW: Understand the workings of the EXPLAIN statement in PostgreSQL

```
# explain select * from people, managers where people.playerid=managers.playerid;  
                                QUERY PLAN  
-----  
Merge Join (cost=265.09..1341.11 rows=3370 width=185)  
Merge Cond: ((people.playerid)::text = (managers.playerid)::text)  
-> Index Scan using people_pkey on people (cost=0.00..982.40 rows=18599 width=144)  
-> Sort (cost=263.16..271.58 rows=3370 width=41)  
    Sort Key: managers.playerid  
    -> Seq Scan on managers (cost=0.00..65.70 rows=3370 width=41)  
(6 rows)
```