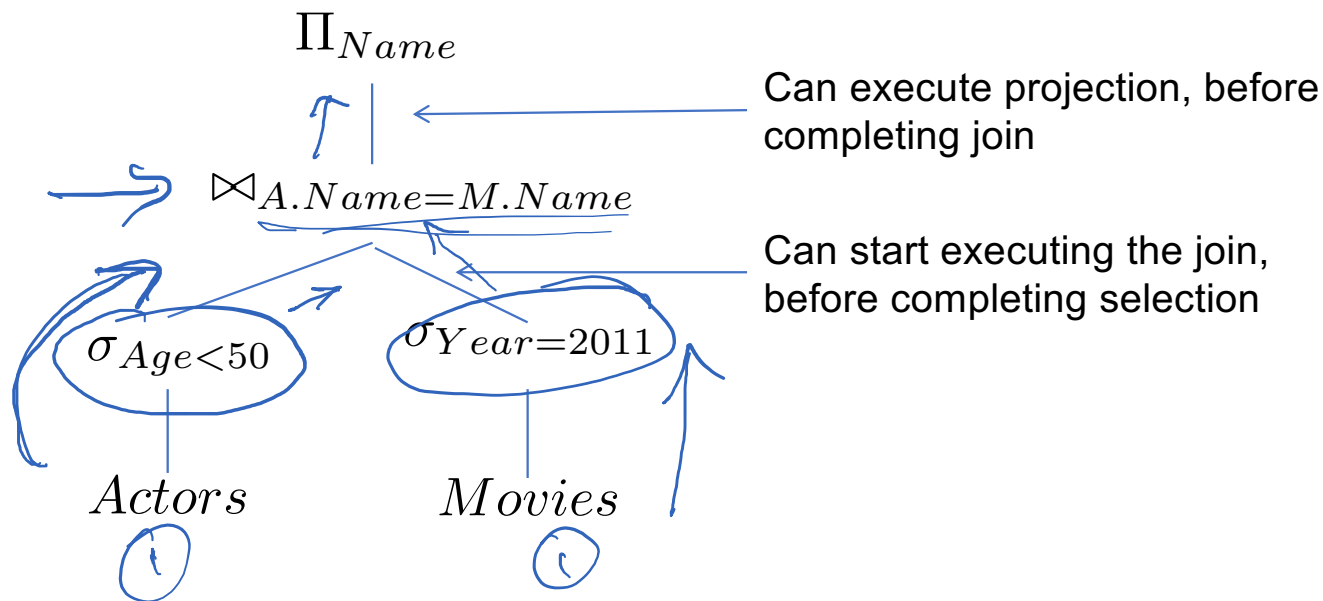


COL 362 & COL 632

Processing SQL – Intro to QOpt, and Storage

15 Feb 2023

Blocking and pipelined execution



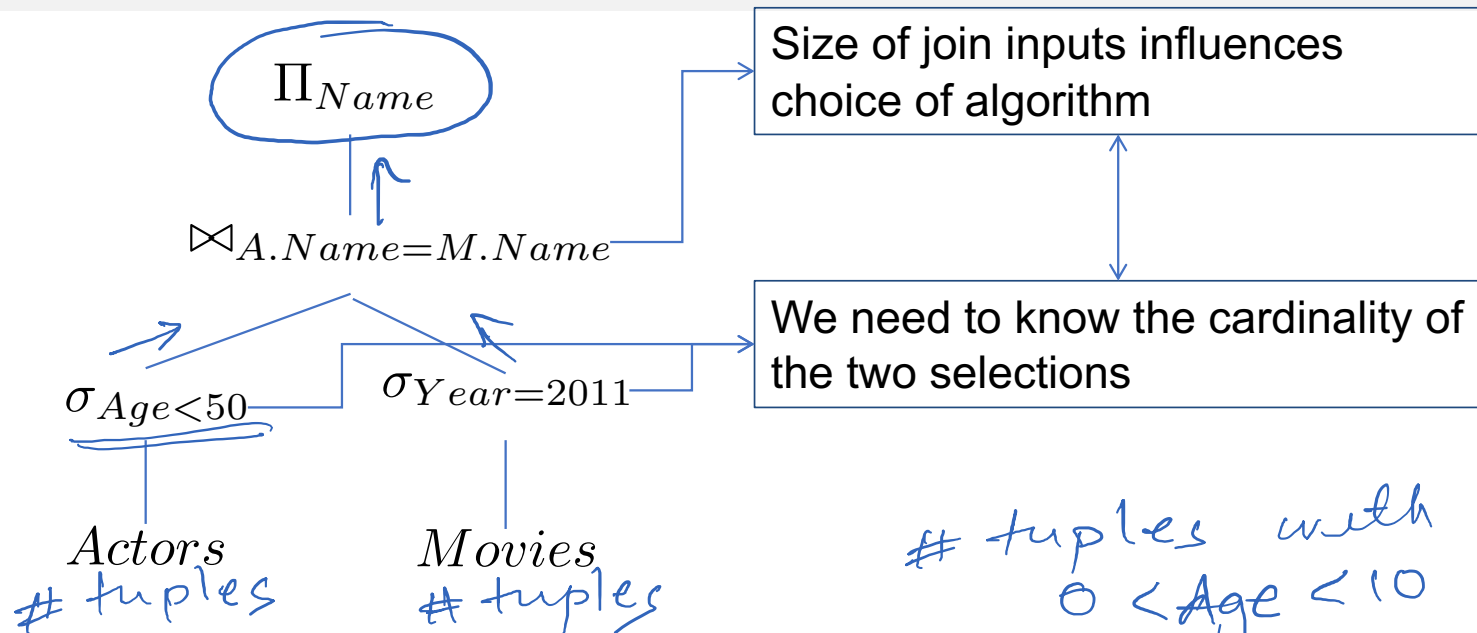
A ←
B
C
:
:
2

What operators are blocking?

Choosing the right operator

- Materialize or pipeline? ✓
- Sort? Index-based? Hash-based? ✓
- Depends on the size of (intermediate) relation
 - Determines amount of main memory required
- “What is the next operator?”
 - Use sort-based, if there is an ORDER BY ahead

Cardinality Estimation (1/2)



Cardinality Estimation (2/2)

- Order of joins matters a lot

$$R \bowtie S \bowtie T$$

- Is it better to join R and S, then T? Or R and T, then S? Or...?
 - Depends on the cardinality
 - Also on other things...

Estimating Cardinalities

- Collect and store statistics about the data
 - What to collect?
 - Statistical summaries should be very small
- Estimate cardinalities for various operators
 - How to use the stats to estimate the cardinalities?
- Up-to-date statistics
 - Accurate statistics lead to accurate estimates

How can you force PostgreSQL to collect statistics?

Statistics in the system

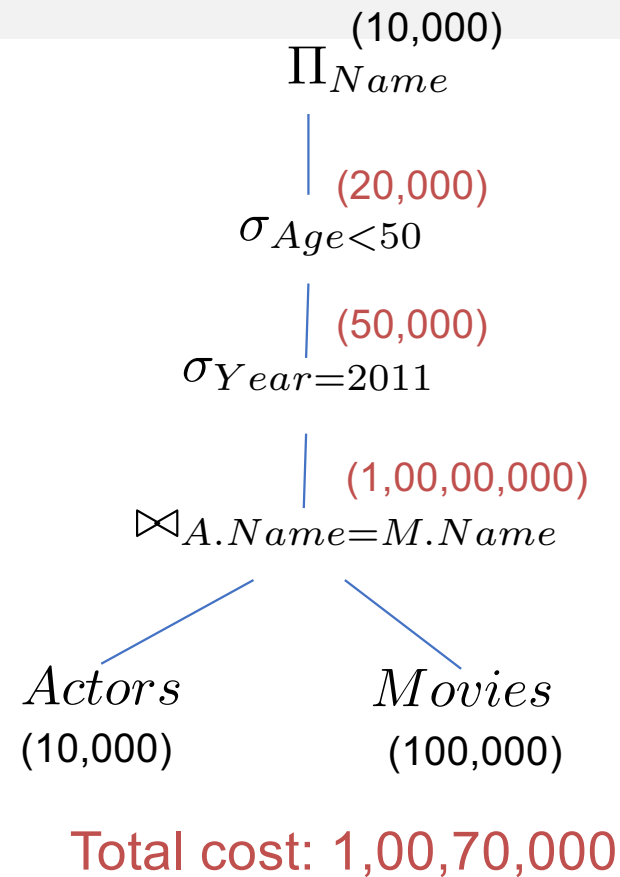
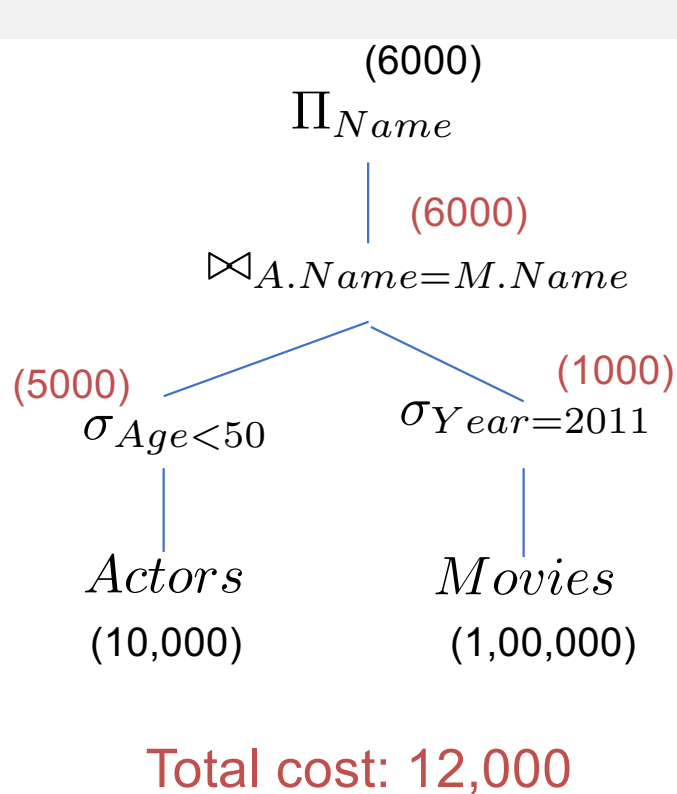
- For table T
 - total number of tuples ✓
 - no. of disk blocks
- For a given attribute
 - Min/max values, no. of distinct values
 - Histogram of values

width of a tuple (# of bytes)

Pg-statistic
Vacuum

What statistics does PostgreSQL collect?

Comparing query plans

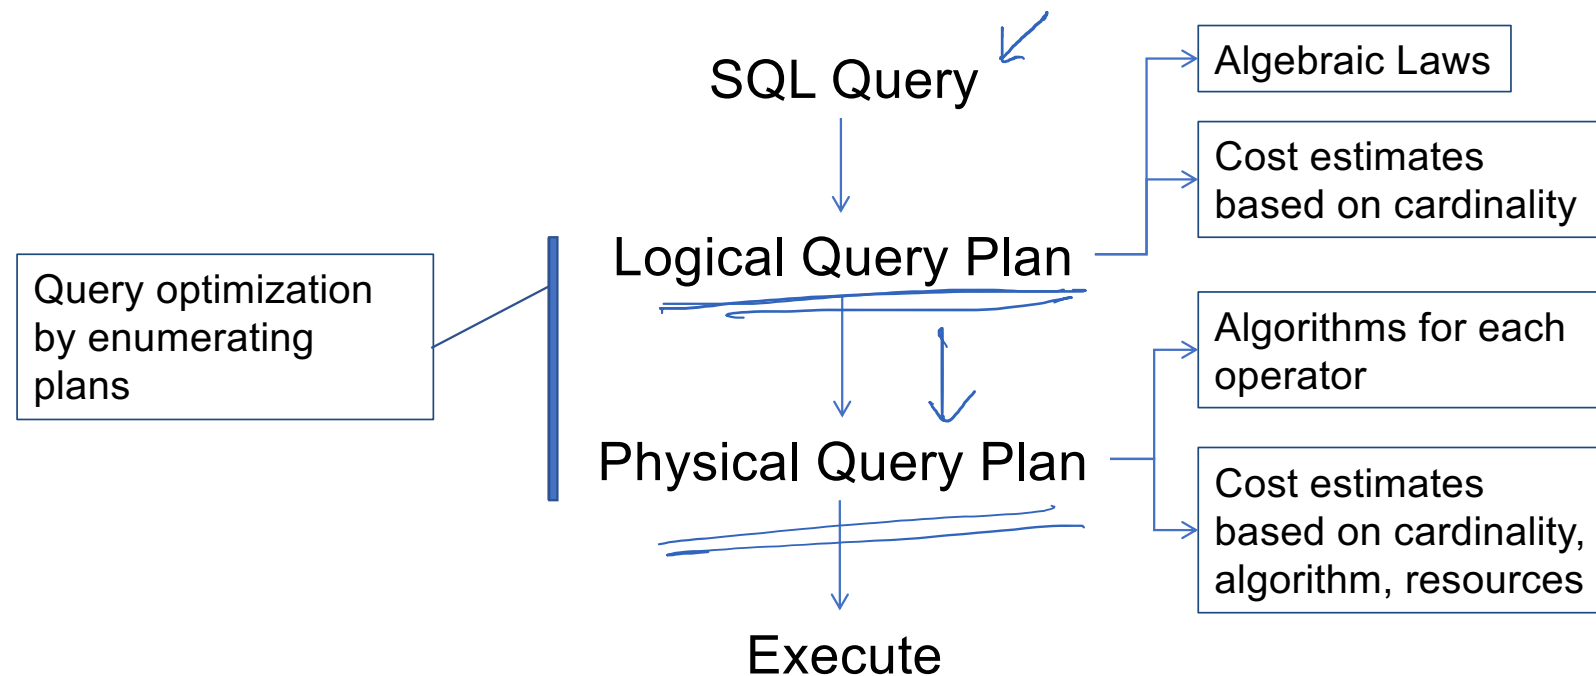


Query plan enumeration

- Top-down
 - Start from the root
 - For each implementation of the root operator, consider all implementations of its children
- Bottom-up
 - For each sub-expression, compute costs of evaluating in all possible ways
 - Dynamic programming

System R
↕
Volcano

Summary

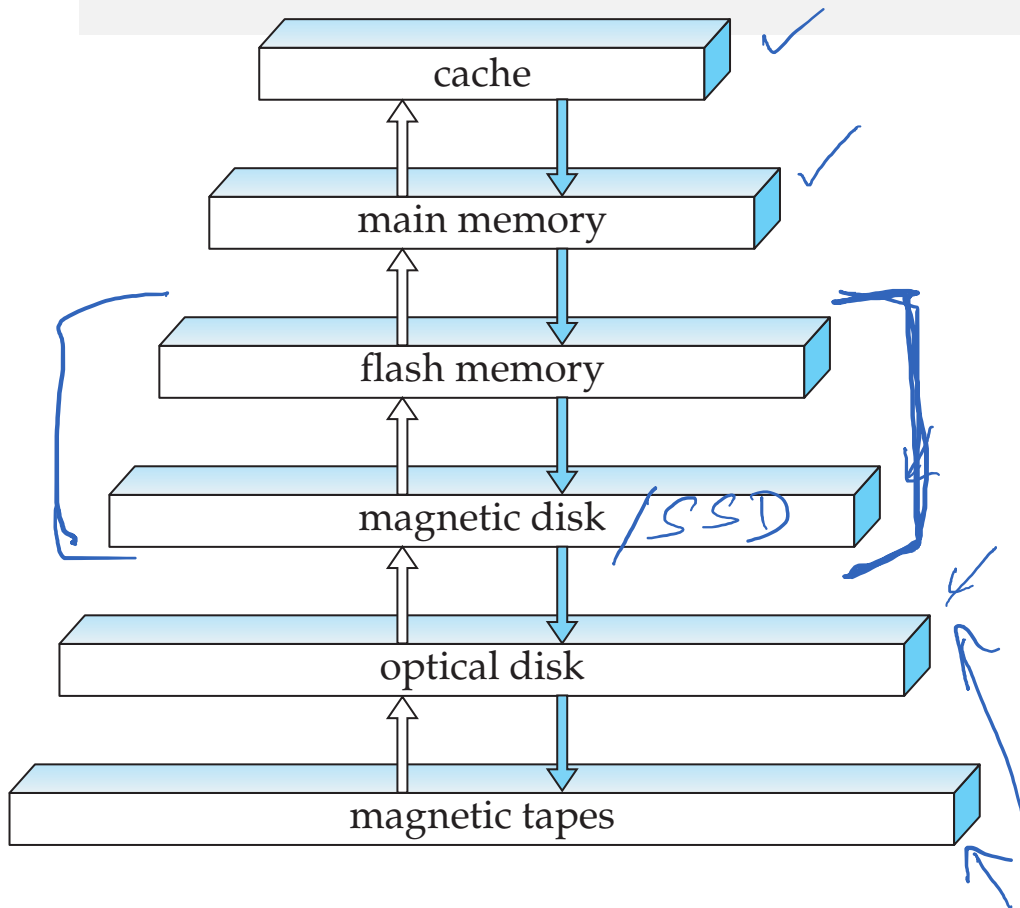


Storage

Classification of Physical Storage Media

- Speed with which data can be accessed
- Cost per unit of data
- Reliability ✓
 - data loss on power failure or system crash
 - physical failure of the storage device
- Can differentiate storage into:
 - **volatile storage**: loses contents when power is switched off
 - **non-volatile storage**:
 - Contents persist even when power is switched off.
 - Includes secondary and tertiary storage, as well as battery-backed up main-memory.

Storage Hierarchy



- Cache
 - Fastest, costliest
- Main Memory
- Flash
 - NAND and NOR flash
 - Used in phones, USB sticks, as SSDs
- Magnetic Disk
- Optical Storage
 - CDs, DVDs
- Tape storage