


COL 362 & COL 632

Database Storage

21 Feb 2023

MTTF


U genix CMU






- MTTF = Mean Time To Failure
- If you had n disks, and you operated them for T hours before f of them failed then $T \cdot n / f$
- *If you test a sample of 1,000 drives for a period of 1,000 hours and then one of them failed, then*

$$MTTF = \frac{1,000 \times 1,000}{1} = 10,00,000 \text{ hours} \sim 114 \text{ years}$$


Now, if you had 114 drives, it is likely that at least one of them will fail after 1 year



Optimization of Disk-Block Access

- Buffering 
 - Fetching blocks from disk to satisfy future requests
 - DBMS implement their own buffering policies 
- Read-ahead 
 - Read extra blocks from a track in anticipation that they will be requested soon
- Disk-arm-scheduling algorithms re-order block requests so that disk arm movement is minimized
 - elevator algorithm

Buffer manager

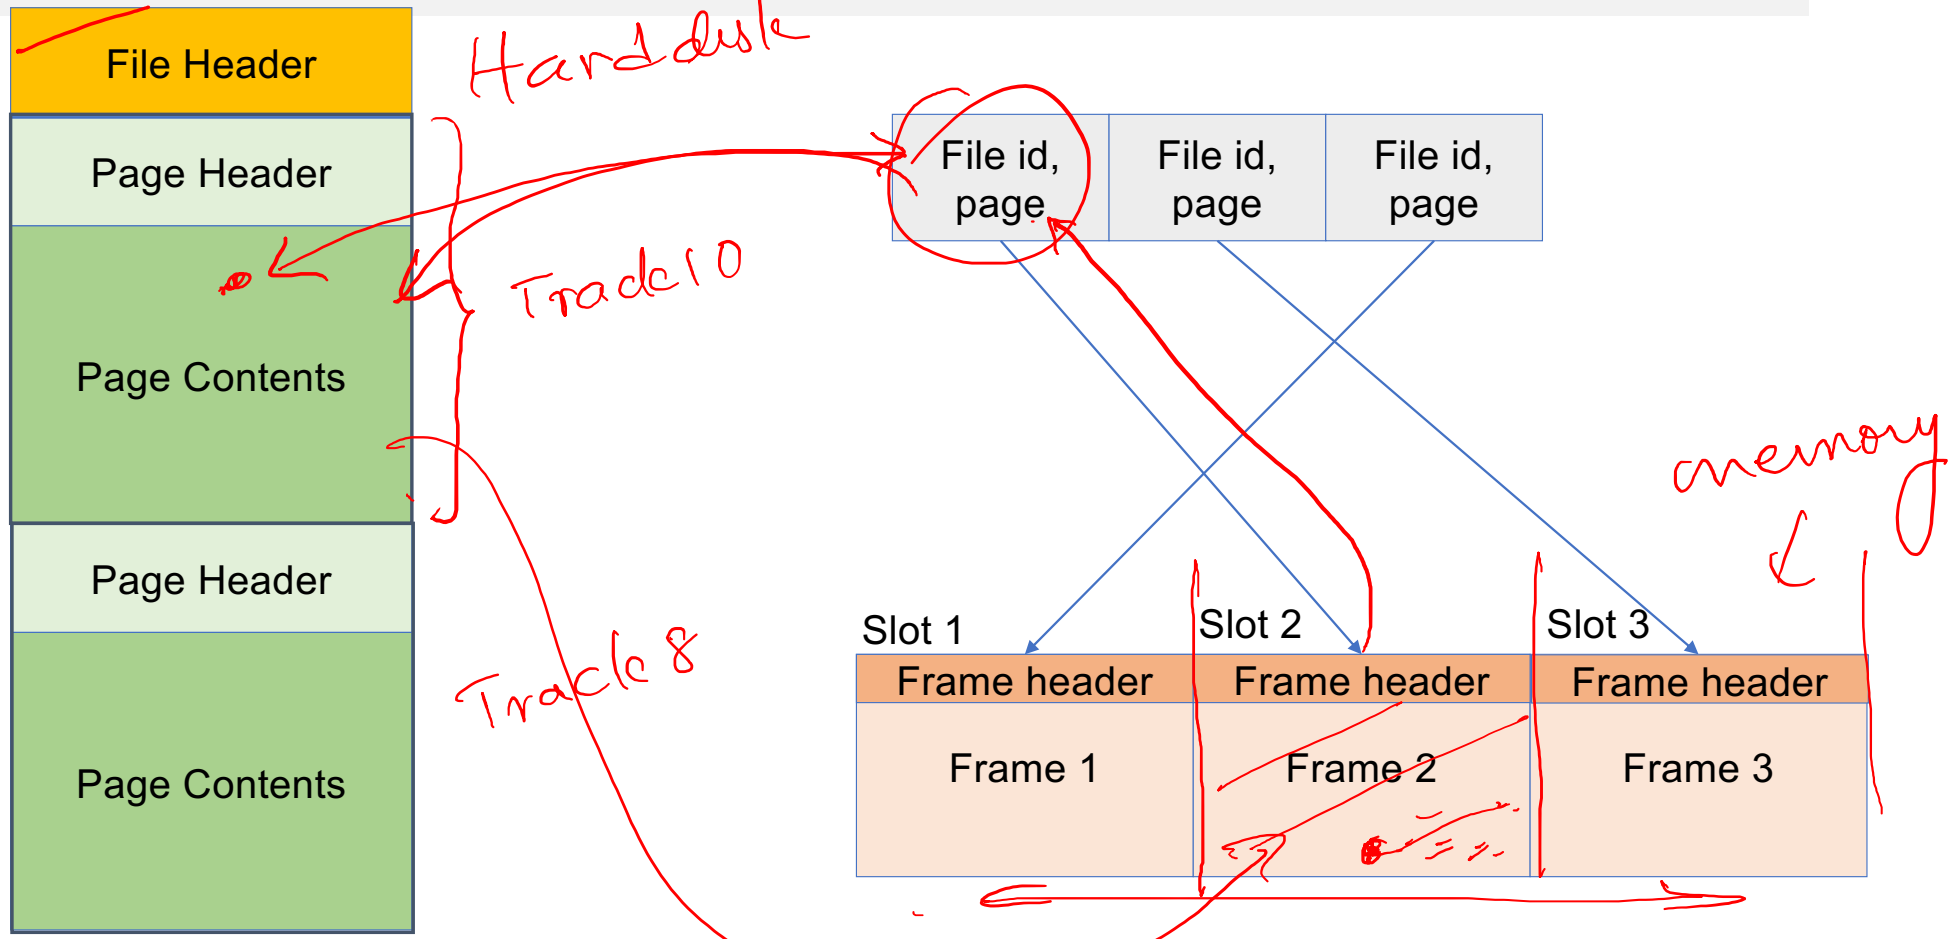


- File Organization
 - Organize files in a way that corresponds to how it will be accessed
 - Ex: contents of a file organized in consecutive blocks

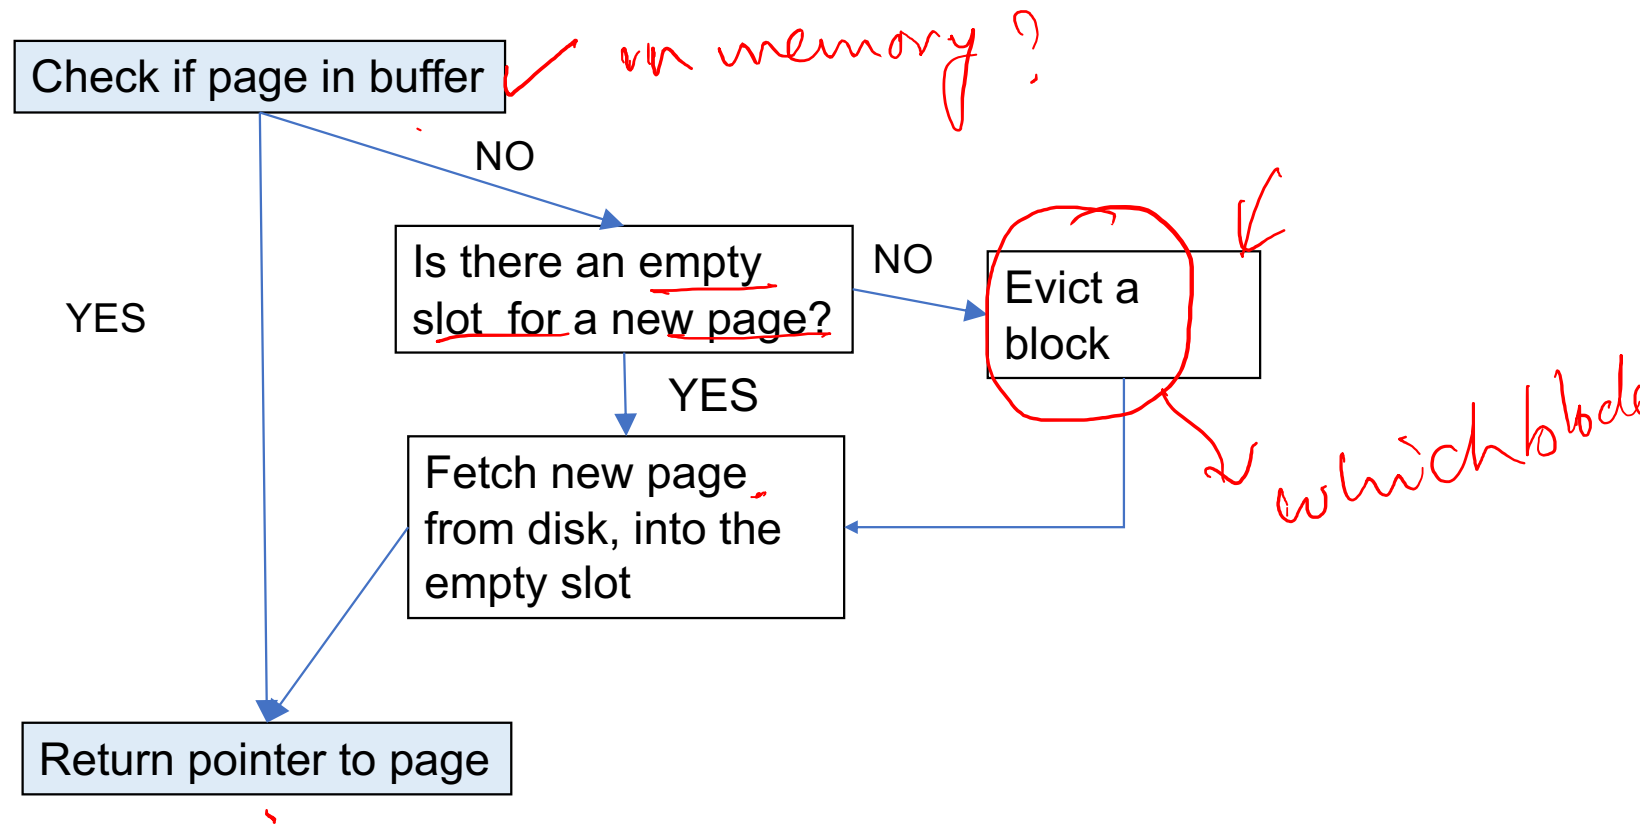
Storage Access

- **Blocks** are units of both storage allocation and data transfer.
 - Goal: minimize transfers from disk to memory
- **Page** is a unit of storage on disk that is structured
- **Buffer** – portion of main memory available to store copies of disk ~~blocks~~ *pages*
- **Frame** is a unit of storage in memory that is structured
- **Buffer manager** – subsystem responsible for allocating buffer ~~space in main memory~~



Files and Buffers



Buffer Manager



Buffer Replacement Policies (1/2)

- Pinned blocks 
 - Blocks which cannot be evicted until unpinned
- Forced output 
 - Flush the block to disk
- Buffer replacement strategy
 - LRU, Clock, MRU, FIFO, etc.

LRU-k



Buffer Replacement Policies (2/2)

B1, B2, B3, B1, B4, B3, B2

	M1	M2	M3
B1	<u>B1</u>	-	-
B2	B1	<u>B2</u>	-
B3	B1	B2	<u>B3</u>
B1	B1	B2	B3
B4	B1	B4	B3
B3	B1	B4	B3
B2	B2	B4	B3

LRU – Least Recently Used

Nested-loop join

```

for each tuple i of instructor do
  for each tuple d of department do
    if i[dept_name] == d[dept_name] then
      <generate join tuple>
    end
  end
end
    
```

- For department blocks B1, B2, B3, B4 (and loop over for next block of instructor tuples)

LRU

	M1	M2	M3
B1	B1	-	-
B2	B1	B2	-
B3	B1	B2	B3
B4	B4	B2	B3

MRU – Most Recently Used

	M1	M2	M3
B1	B1	-	-
B2	B1	B2	-
B3	B1	B2	B3
B4	B1	B2	B4