

COL 362 & COL 632

Indexing
22 Feb 2023

Basic Concepts

file \equiv relation

- Indexing mechanisms speed up access to desired data.
- **Search Key** – set of attributes used to look up records in a **file**.
- An **index file** consists of records (called **index entries**) of the form



- Index files are typically much smaller than the original file
- Two basic kinds of indexes
 - ✓ **Ordered indexes:** search keys are stored in sorted order
 - **Hash indexes:** search keys are distributed uniformly across buckets using a hash function

Secondary memory Indexes

Dimensions for Evaluating Indexes

- **Access type**

- Tuple with a specified value in the attribute ✓
- Tuples with an attribute value falling in a specified range of values. ✓

Hash indexes

- ✓ • **Access time**

- Find one or more items – or to confirm searched item does not exist

- **Insertion time**

- **Deletion time**

- **Space overhead** ✓

More than one index for a relation

Ordered Indexes

- Index entries are sorted on the **search key** value
 - Search key can be **one or more** attributes
- Primary or clustering index
 - Search key is the same as the corresponding sequentially ordered file
 - The search key of a primary index is usually *but not necessarily* the primary key
- Secondary or non-clustering index
 - Search key different from corresponding sequentially ordered file

Dense Indexes (1/2)

Primary key is a
Candidate key

- Index record appears for every search-key value in the file.
- index on ID attribute of instructor relation

10101
10101

Index	File	Name				
10101	10101	Srinivasan	Comp. Sci.	65000		
12121	12121	Wu	Finance	90000		
15151	15151	Mozart	Music	40000		
22222	22222	Einstein	Physics	95000		
32343	32343	El Said	History	60000		
33456	33456	Gold	Physics	87000		
45565	45565	Katz	Comp. Sci.	75000		
58583	58583	Califieri	History	62000		
76543	76543	Singh	Finance	80000		
76766	76766	Crick	Biology	72000		
83821	83821	Brandt	Comp. Sci.	92000		
98345	98345	Kim	Elec. Eng.	80000		

Dense Indexes (2/2)

Dense clustered Index
Dense non-clustered Index

- Dense index on dept_name, with instructor file sorted on dept_name

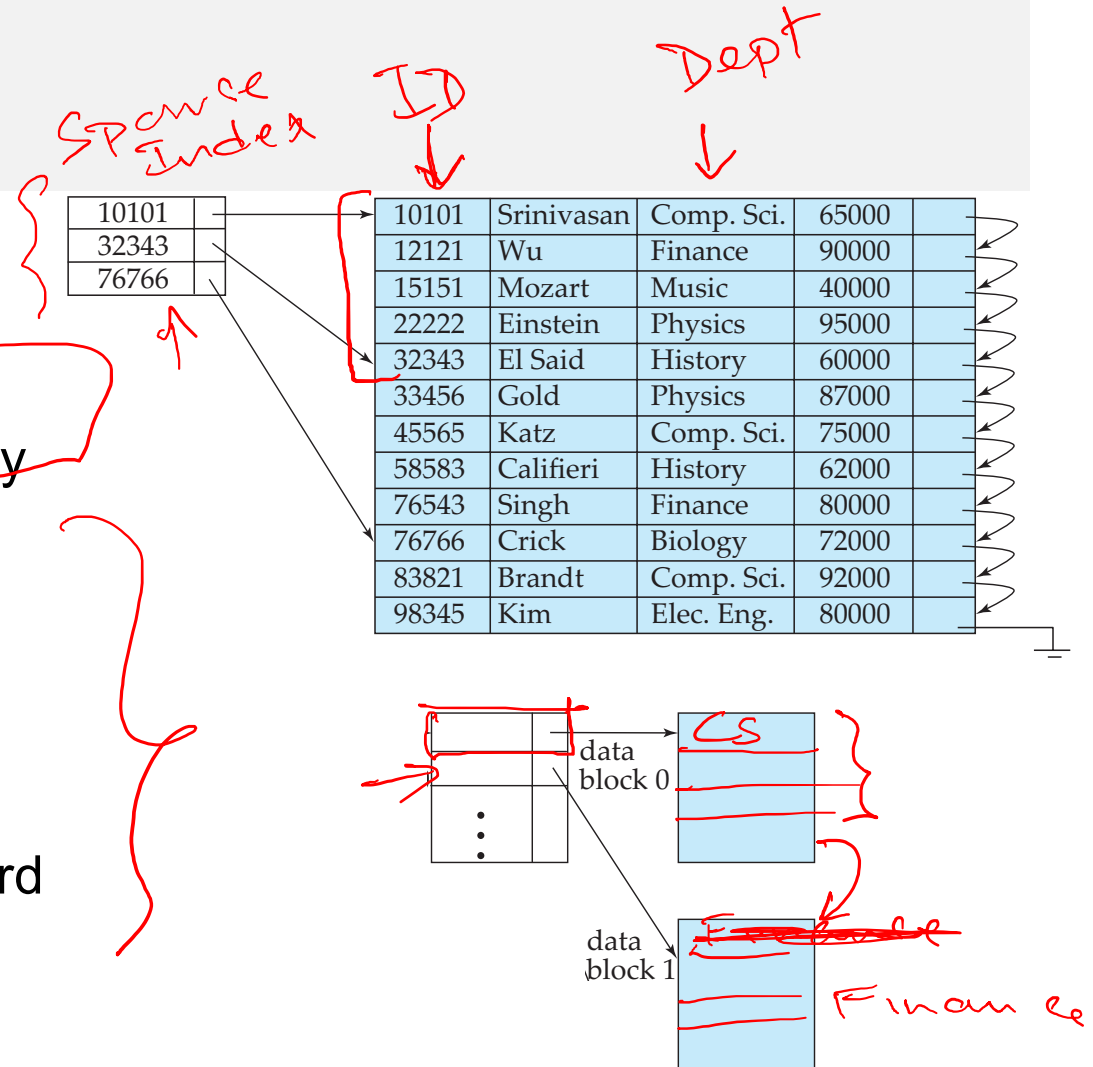
Biology	76766	Crick	Biology	72000
Comp. Sci.	10101	Srinivasan	Comp. Sci.	65000
Elec. Eng.	45565	Katz	Comp. Sci.	75000
Finance	83821	Brandt	Comp. Sci.	92000
History	98345	Kim	Elec. Eng.	80000
Music	12121	Wu	Finance	90000
Physics	76543	Singh	Finance	80000
	32343	El Said	History	60000
	58583	Califieri	History	62000
	15151	Mozart	Music	40000
	22222	Einstein	Physics	95000
	33465	Gold	Physics	87000

Dense

Biology
Finance
Physics

Sparse Indexes


- Contains index records for only some search-key values
 - Applicable when records are sequentially ordered on search-key
- To locate a record with search-key value K we:
 - Find index record with largest search-key value $< K$
 - Search file sequentially starting at the record to which the index record points



Multilevel Indexes

- If primary index does not fit in memory, access becomes expensive.
- Construct a sparse index on primary index – **2-level index**
- Extend the idea to multiple levels

B⁺-tree




Insertion

Dense


Biology		76766	Crick	Biology	72000	
Comp. Sci.		10101	Srinivasan	Comp. Sci.	65000	
Elec. Eng.		45565	Katz	Comp. Sci.	75000	
Finance		83821	Brandt	Comp. Sci.	92000	
History		98345	Kim	Elec. Eng.	80000	
Music		12121	Wu	Finance	90000	
Physics		76543	Singh	Finance	80000	
		32343	El Said	History	60000	
		58583	Califieri	History	62000	
		15151	Mozart	Music	40000	
		22222	Einstein	Physics	95000	
		33465	Gold	Physics	87000	

Sparse



Deletion

Biology	76766	Crick	Biology	72000
Comp. Sci.	10101	Srinivasan	Comp. Sci.	65000
Elec. Eng.	45565	Katz	Comp. Sci.	75000
Finance	83821	Brandt	Comp. Sci.	92000
History	98345	Kim	Elec. Eng.	80000
Music	12121	Wu	Finance	90000
Physics	76543	Singh	Finance	80000
	32343	El Said	History	60000
	58583	Califieri	History	62000
	15151	Mozart	Music	40000
	22222	Einstein	Physics	95000
	33465	Gold	Physics	87000



Secondary Indexes

- Index record points to a bucket that contains pointers to all the actual records with that particular search-key value.
- Secondary indexes have to be dense

