

*Indexed
Sequential
Files*

COL 362 & COL 632

Indexing

01 Mar 2023

Insertion

| | | | | | | |
|------------|--|-------|------------|------------|-------|--|
| Biology | | 76766 | Crick | Biology | 72000 | |
| Comp. Sci. | | 10101 | Srinivasan | Comp. Sci. | 65000 | |
| Elec. Eng. | | 45565 | Katz | Comp. Sci. | 75000 | |
| Finance | | 83821 | Brandt | Comp. Sci. | 92000 | |
| History | | 98345 | Kim | Elec. Eng. | 80000 | |
| Music | | 12121 | Wu | Finance | 90000 | |
| Physics | | 76543 | Singh | Finance | 80000 | |
| | | 32343 | El Said | History | 60000 | |
| | | 58583 | Califieri | History | 62000 | |
| | | 15151 | Mozart | Music | 40000 | |
| | | 22222 | Einstein | Physics | 95000 | |
| | | 33465 | Gold | Physics | 87000 | |


Dense

Energy Sciences

Overflow

The diagram illustrates a relationship between two data structures. On the left is a table with 'Energy' in its first row. On the right is a larger table with 'Energy' in its first row and 'Material science' in its second row. A line connects the 'Energy' cell of the left table to the 'Energy' cell of the right table.

Spartle




Sparse index

- ① Space is in the page
- ② create new page

Deletion

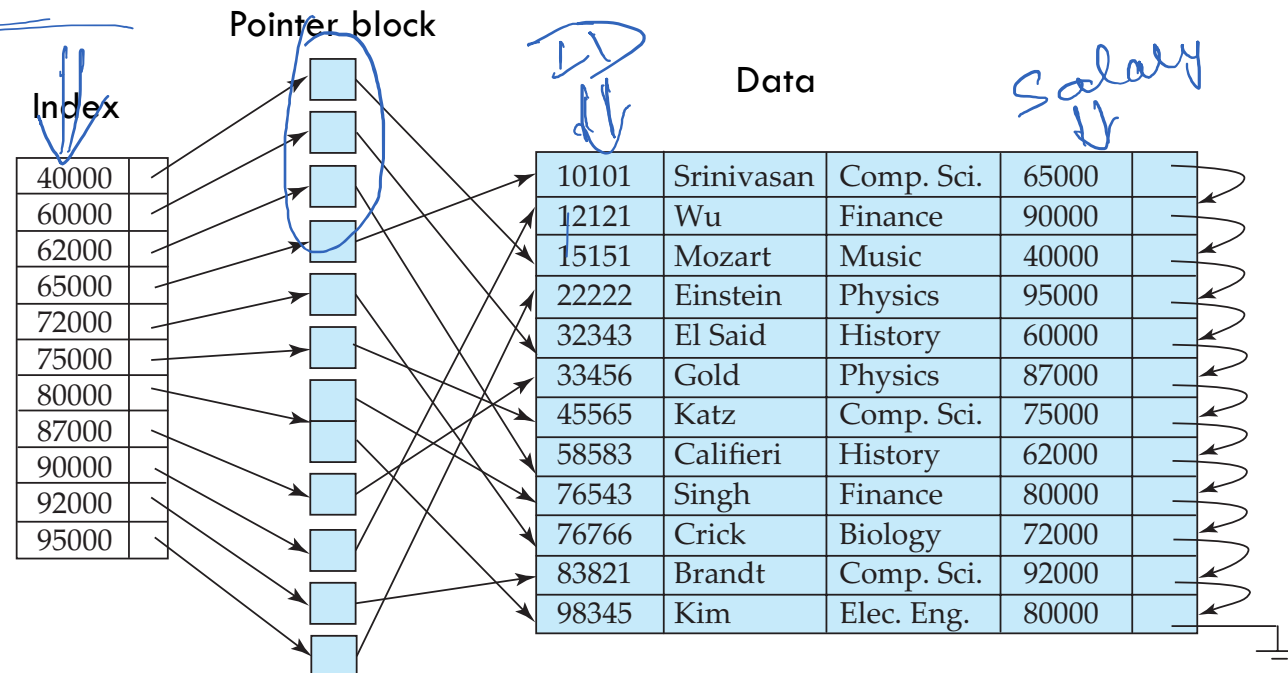
H. W.

| | | | | | | |
|------------|--|-------|------------|------------|-------|--|
| Biology | | 76766 | Crick | Biology | 72000 | |
| Comp. Sci. | | 10101 | Srinivasan | Comp. Sci. | 65000 | |
| Elec. Eng. | | 45565 | Katz | Comp. Sci. | 75000 | |
| Finance | | 83821 | Brandt | Comp. Sci. | 92000 | |
| History | | 98345 | Kim | Elec. Eng. | 80000 | |
| Music | | 12121 | Wu | Finance | 90000 | |
| Physics | | 76543 | Singh | Finance | 80000 | |
| | | 32343 | El Said | History | 60000 | |
| | | 58583 | Califieri | History | 62000 | |
| | | 15151 | Mozart | Music | 40000 | |
| | | 22222 | Einstein | Physics | 95000 | |
| | | 33465 | Gold | Physics | 87000 | |



Secondary Indexes

- Index record points to a bucket that contains pointers to all the actual records with that particular search-key value.
- Secondary indexes have to be dense



|| Goetz Graefe's
monograph


B+-trees

B⁺-Trees (1/2)

Ubiquitous

(2,3) - trees + + +
(a, b)

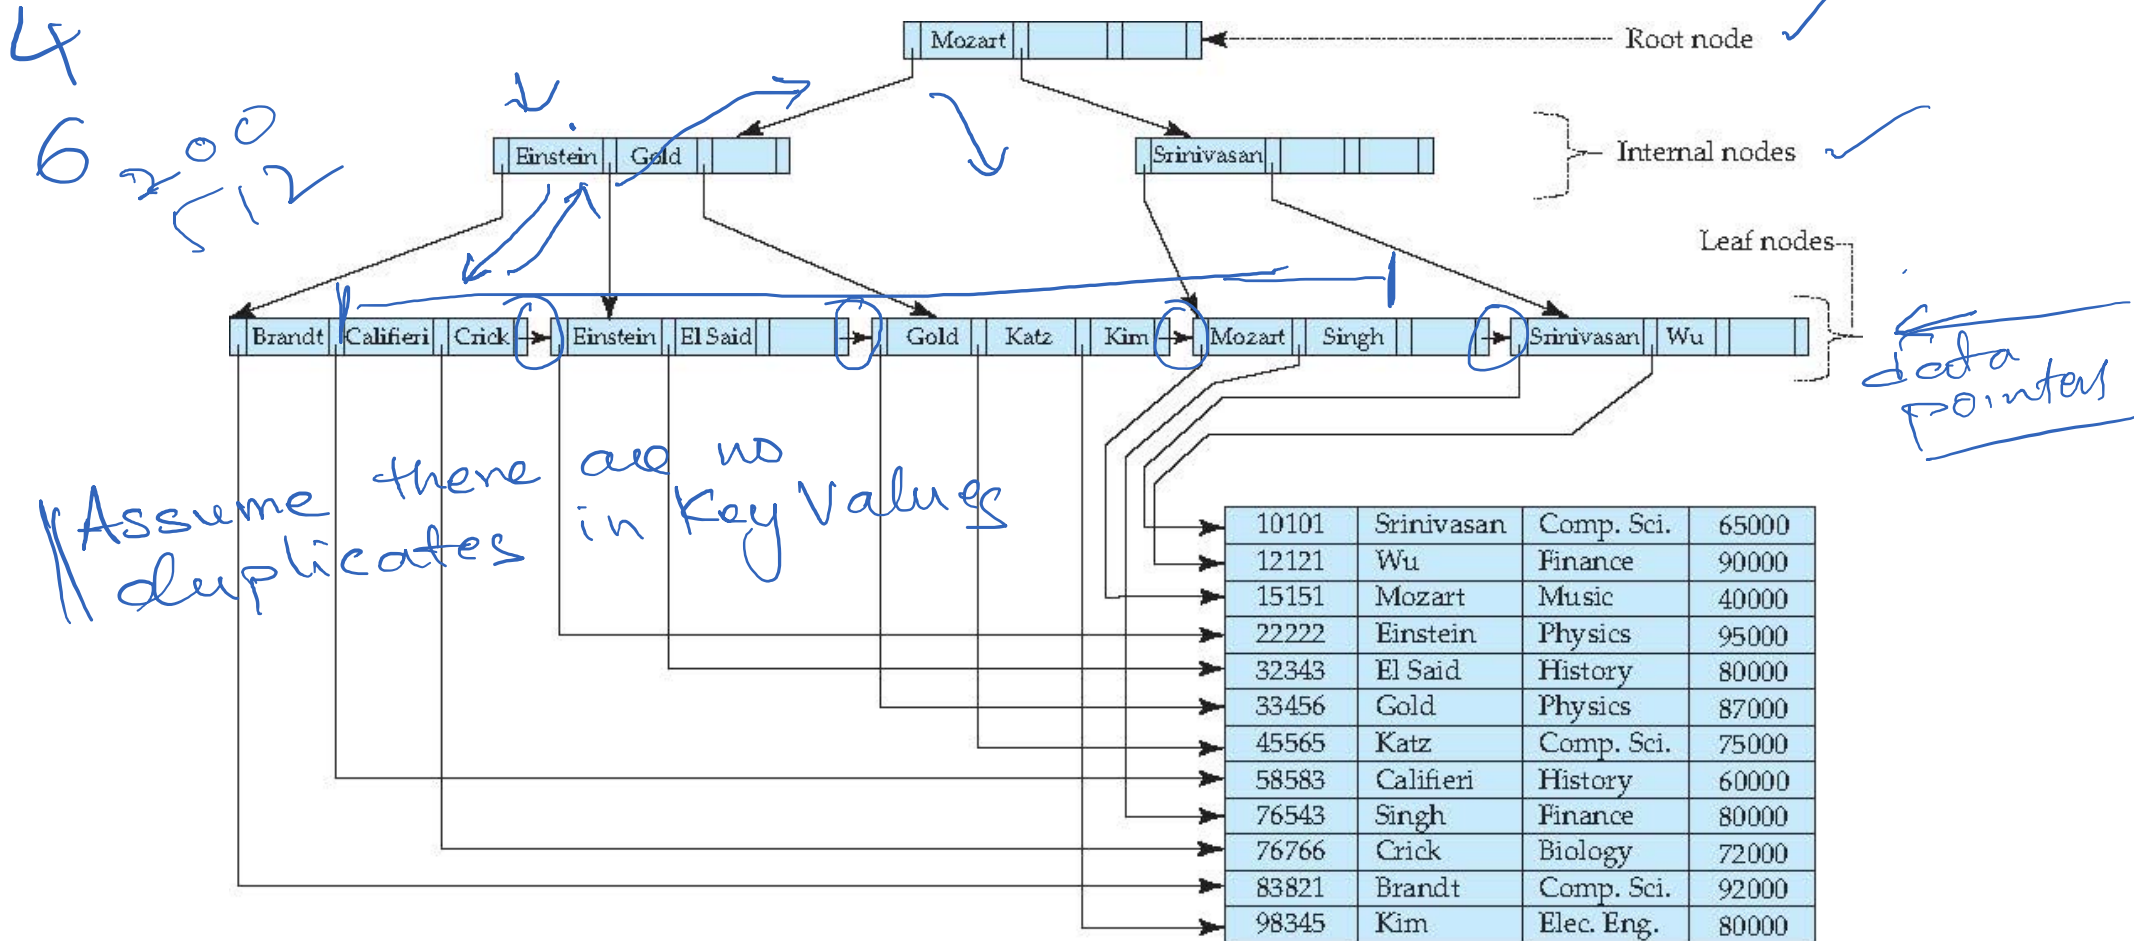
- Disadvantage of indexed-sequential files
 - performance degrades as file grows, since many overflow blocks get created
 - Periodic reorganization of entire file is required
- Advantage of B⁺-tree index files
 - automatically reorganizes itself
 - Reorganization of entire file is not required
- (Minor) disadvantage of B⁺-trees
 - extra insertion and deletion overhead, space overhead




B+-Tree (2/2)

$$\log_{\lceil n/2 \rceil} N$$

$$\lceil \frac{n}{2} \rceil \leq \# \text{ of children} \leq n$$



B⁺-Tree Node Structure



- K_i are the search-key values
- P_i are pointers to children (for non-leaf nodes) or pointers to records or buckets of records (for leaf nodes).
- The search-keys in a node are ordered

$$K_1 < K_2 < K_3 < \dots < K_{n-1}$$

(Assuming no duplicate keys)

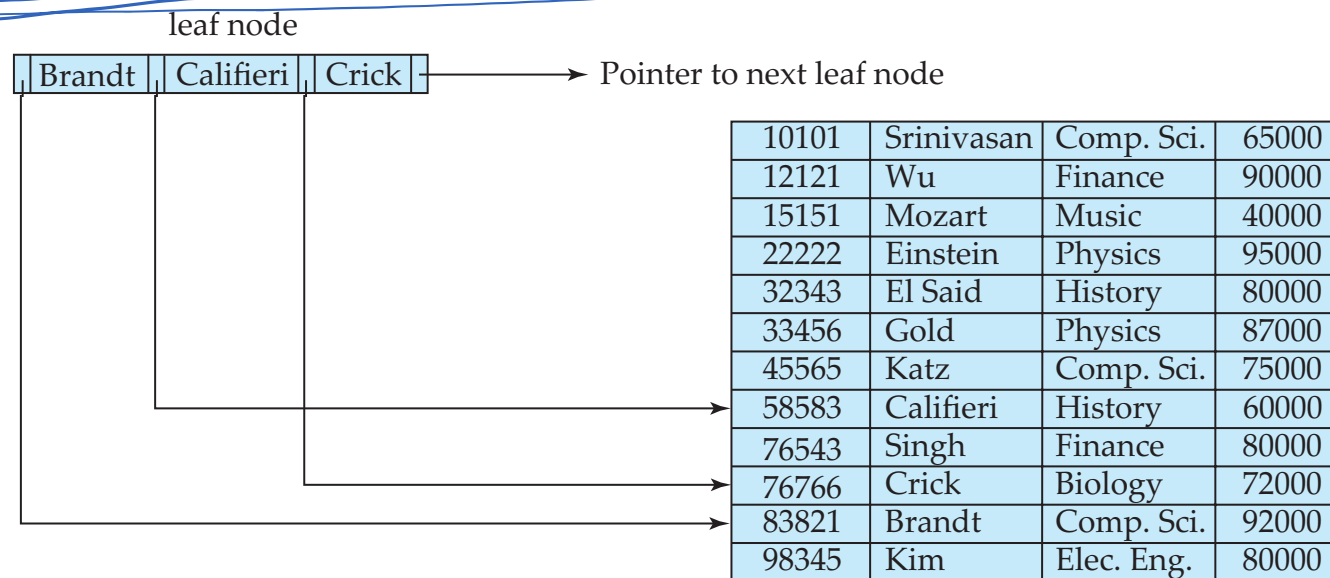
Block-linked
list

Leaf Nodes in B⁺-Trees

look like
dense index

- For $i = 1, 2, \dots, n-1$, pointer P_i points to a file record with search-key value K_i ,
- If L_i, L_j are leaf nodes and $i < j$, L_i 's search-key values are less than or equal to L_j 's search-key values
- P_n points to next leaf node in search-key order


P_n




Non-Leaf Nodes in B⁺-Trees

$\lceil n/2 \rceil$ $[n]$

- For a non-leaf node with m pointers
 - All the search-keys in the subtree to which P_1 points are less than K_1
 - For $2 \leq i \leq n - 1$, all the search-keys in the subtree to which P_i points have values greater than or equal to K_{i-1} and less than K_i
 - All the search-keys in the subtree to which P_n points have values greater than or equal to K_{n-1}



Example of B⁺-tree




B⁺-tree for `instructor` file ($n = 6$)

- Leaf nodes must have between 3 and 5 **values** ($\lceil (n-1)/2 \rceil$ and $n-1$, with $n = 6$).
- Non-leaf nodes other than root must have between 3 and 6 **children** ($\lceil n/2 \rceil$ and n with $n = 6$).
- Root must have at least 2 children.

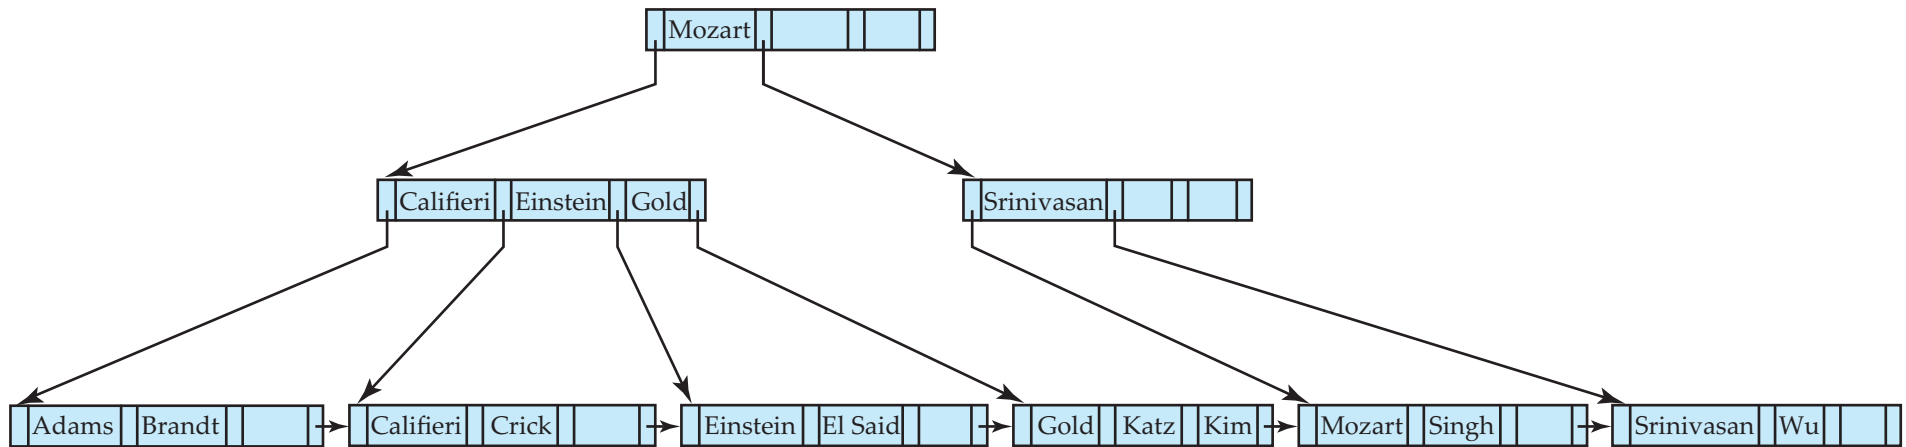
$n = 6$
6

Observations about B⁺-trees

- Since the inter-node connections are done by pointers, “logically” close blocks need not be “physically” close.
- The non-leaf levels of the B⁺-tree form a hierarchy of sparse indices.
- The B⁺-tree contains a relatively small number of levels
 - “Short and fat”
- Insertions and deletions are done in logarithmic time



Queries on B⁺-Trees



Find record with search-key value V .

1. $C = \text{root}$
2. While C is not a leaf node {
 1. Let i be least value s.t. $V \leq K_i$.
 2. If no such exists, set $C = \text{last non-null pointer in } C$
 3. Else { if $(V = K_i)$ Set $C = P_{i+1}$ else set $C = P_i$ }
3. Let i be least value s.t. $K_i = V$
4. If there is such a value i , follow pointer P_i to the desired record.
5. Else no record with search-key value V exists.

Analysis of B⁺-trees

- No. of search keys: 1,000,000
- Block size: 4K
- Size of an index entry: 40B
- Max. no. of search keys per block:
- Max. height of the tree:
- No. of block accesses per query: