


# **COL 362 & COL 632**

Return of Query Processing

15 Mar 2023

# Recap



- Utilize logical plan equivalences to rewrite queries to those which are more efficient
- Choose the most efficient physical operator
- Order the overall query plan to minimize the cost
- For simplicity we just use the number of block transfers from disk and the number of seeks as the cost measure
  - $t_T$  – time to transfer one block
  - $t_S$  – time for one seek
  - Cost for  $b$  block transfers plus  $S$  seeks:  $b * t_T + S * t_S$
- We ignore CPU costs

# Measure of Query Cost – in reality

- Required data may be buffer resident already, avoiding disk I/O
  - But hard to take into account for cost estimation
- Several algorithms can reduce disk IO by using extra buffer space
  - Amount of real memory available to buffer depends on other concurrent queries and OS processes, known only during execution
- **Worst case** estimates assume that no data is initially in buffer and only the minimum amount of memory needed for the operation is available
  - But more optimistic estimates are used in practice

# Selection Operation

$$\text{linear} = b_r^{\text{read}} + \text{seek}$$

$$\text{binary} = \log_2 b_r \times \text{read} + \log_2 b_r \times \text{seek}$$

- **File scan**

- Algorithm **A1** (**linear search**). Scan each file block and test all records to see whether they satisfy the selection condition.

- Cost estimate =  $b_r$  block transfers + 1 seek

- $b_r$  denotes number of blocks containing records from relation  $r$

- If selection is on a key attribute, can stop on finding record }



- cost =  $(b_r/2)$  block transfers + 1 seek

selection attribute

- Can we use binary search?

→ sorted on the selection attribute  
→ seek count goes up

# Selections Using Indices

$$\sigma_k = v$$

- **Index scan** – search algorithms that use an index
  - selection condition must be on search-key of index.
- **A2 (clustering index, equality on key)**. Retrieve a single record that satisfies the corresponding equality condition
  - $\text{Cost} = (h_i + 1) * (t_T + t_S)$  ←
- **A3 (clustering index, equality on nonkey)** Retrieve multiple records.

B+ tree  
Hash

No overflow  
buckets

- Records will be on consecutive blocks  
Let  $b$  = number of blocks containing matching records

- $\text{Cost} = h_i * (t_T + t_S) + t_S + t_T * b$

$$(h_i + 1) (t_S + t_T) + (b - 1) \times t_T$$

# Selections Using Indices

- **A4 (secondary index, equality on key/non-key).**
  - Retrieve a single record if the search-key is a candidate key  
 $\text{Cost} = (h_i + 1) * (t_T + t_S)$
  - Retrieve multiple records if search-key is not a candidate key  
=> each of  $n$  matching records may be on a different block  
 $\text{Cost} = (h_i + n) * (t_T + t_S)$ 
    - Can be very expensive!

# Selections Involving Comparisons

- Can implement selections of the form  $\sigma_{A \leq v}(r)$  or  $\sigma_{A \geq v}(r)$  by using
  - a linear file scan, ✓
  - or by using indices in the following ways:
- **A5 (clustering index, comparison).** (Relation is sorted on A)
  - For  $\sigma_{A \geq v}(r)$  use index to find first tuple  $\geq v$  and scan relation sequentially from there
  - For  $\sigma_{A \leq v}(r)$  just scan relation sequentially till first tuple  $> v$ ; do not use index
- **A6 (clustering index, comparison).**
  - For  $\sigma_{A \geq v}(r)$  use index to find first index entry  $\geq v$  and scan index sequentially from there, to find pointers to records.
  - For  $\sigma_{A \leq v}(r)$  just scan leaf pages of index finding pointers to records, till first entry  $> v$ .
  - In either case, retrieve records that are pointed to
  - requires an I/O per record; Linear file scan may be cheaper!

$$\theta_1 \wedge \theta_2 \equiv \theta_2 \wedge \theta_1$$

# Implementation of Complex Selections

$$\frac{A_1 > V_1}{\theta_1} \wedge \frac{A_2 = V_2}{\theta_2}$$

- **Conjunction:**  $\sigma_{\theta_1 \wedge \theta_2 \wedge \dots \wedge \theta_n}(r)$
- **A7 and A8**
- **A9 (conjunctive selection by intersection of identifiers).**
  - Requires indices with record pointers.
  - Use corresponding index for each condition, and take intersection of all the obtained sets of record pointers.
  - Then fetch records from file
  - If some conditions do not have appropriate indices, apply test in memory.



HW

# Algorithms for Complex Selections

- **Disjunction:**  $\sigma_{\theta_1 \vee \theta_2 \vee \dots \vee \theta_n}(r)$ .
- **A10 (disjunctive selection by union of identifiers).**
  - Applicable if *all* conditions have available indices.
    - Otherwise use linear scan.
  - Use corresponding index for each condition and take union of all the obtained sets of record pointers.
  - Then fetch records from file
- **Negation:**  $\sigma_{\neg\theta}(r)$ 
  - Use linear scan on file
  - If very few records satisfy  $\neg\theta$ , and an index is applicable to  $\theta$ 
    - Find satisfying records using index and fetch from file