

COL 362 & COL 632

Analysis of Sorting, Joins

21 Mar 2023


Reducing Seeks in External Sort

- 1 block per run leads to too many seeks during merge
- Instead use b_b buffer blocks per run
 \Rightarrow read/write b_b blocks at a time from each run
- Can merge $\left\lfloor \frac{M}{b_b} \right\rfloor - 1$ runs in one pass
- Total number of merge passes required:

$$\left\lceil \log_{\left\lfloor \frac{M}{b_b} \right\rfloor - 1} \frac{b_r}{M} \right\rceil$$

\downarrow b_r blocks
 \uparrow M

b_r blocks
 b_b blocks
 per
 run,
 (buffer)



External Sort Merge – Block Transfers

- Block transfers for initial run creation as well as in each pass is $2b_r$
 - For the final pass, we don't count write cost *when can we ignore the final pass.?*
 - We can ignore final write cost for all operations since the output of an operation may be sent to the parent operation without being written to disk
- Thus, total number of **block transfers** for external sorting: *time*

$$b_r \left(2 \left\lceil \log_{\left\lfloor \frac{M}{b_b} \right\rfloor - 1} \frac{b_r}{M} \right\rceil + 1 \right)$$

Handwritten red version of the formula:

$$2b_r \left\lceil \log_{\left\lfloor \frac{M}{b_b} \right\rfloor - 1} \frac{b_r}{M} \right\rceil + \underline{\underline{b_r}}$$

External Merge Sort - Cost of seeks

- **During run generation:** one seek to read each run and one seek to write each run

- $2 \left\lceil \frac{b_r}{M} \right\rceil$

- **During the merge phase:**

- Need $2 \left\lceil \frac{b_r}{b_b} \right\rceil$ seeks for each merge pass

- except the final one which does not require a write

- **Total number of seeks:**

$$2 \left\lceil \frac{b_r}{M} \right\rceil + \left\lceil \frac{b_r}{b_b} \right\rceil \left(2 \left\lceil \log_{\left\lceil \frac{M}{b_b} \right\rceil} \frac{b_r}{M} \right\rceil - 1 \right)$$

Ramakrishnan, Gehrke . - Cow Book

Number of Passes (Tentative)

N	$b_b = 3$	$b_b = 5$	$b_b = 9$	$b_b = 17$	$b_b = 129$	$b_b = 257$
100	7	4	3	2	1	1
1,000	10	5	4	3	2	2
10,000	13	7	5	4	2	2
100,000	17	9	6	5	3	3
1,000,000	20	10	7	5	3	3
10,000,000	23	12	8	6	4	3
100,000,000	26	14	9	7	4	4
1,000,000,000	30	15	10	8	5	4

Up-to-date sort performance at <http://sortbenchmark.org>

Join Operation

- Several different algorithms to implement joins

physical plans

- Nested-loop join
- Block nested-loop join
- Indexed nested-loop join
- Merge-join
- Hash-join

← extn. from external sort merge

- Choice based on cost estimate

- Examples use the following information

- Number of records of *student*: 5,000 takes: 10,000
- Number of blocks of *student*: 100 takes: 400

Nested-Loop Join

- To compute the theta join $r \bowtie_{\theta} s$
 for each tuple t_r in r do begin
 for each tuple t_s in s do begin
 test pair (t_r, t_s) to see if they satisfy the join condition θ
 if they do, add $t_r \cdot t_s$ to the result.
 end
 end
- r is called the outer relation and s the inner relation of the join.
- Requires no indices and can be used with any kind of join condition.
- Expensive since it examines every pair of tuples in the two relations.

Nested-Loop Join (Cont.)

- In the **worst case**, if there is enough memory only to hold one block of each relation, the estimated cost is $n_r * bs + b_r$ block transfers, plus $n_r + b_r$ seeks
- Assuming worst case memory availability cost estimate is
 - with *student* as outer relation:
 - $5000 * 400 + 100 = 2,000,100$ block transfers,
 - $5000 + 100 = 5100$ seeks
 - with *takes* as the outer relation
 - $10000 * 100 + 400 = 1,000,400$ block transfers and 10,400 seeks
- If the smaller relation fits entirely in memory, use that as the inner relation
 - Reduces cost to $b_r + bs$ block transfers and 2 seeks

Block Nested-Loop Join

- Variant of nested-loop join in which every block of inner relation is paired with every block of outer relation.

```
for each block  $B_r$  of  $r$  do begin
  for each block  $B_s$  of  $s$  do begin
    for each tuple  $t_r$  in  $B_r$  do begin
      for each tuple  $t_s$  in  $B_s$  do begin
        Check if  $(t_r, t_s)$  satisfy the join condition
        if they do, add  $t_r \cdot t_s$  to the result.
      end
    end
  end
end
end
```



Block Nested-Loop Join (Cont.)

- Worst case estimate: $b_r * b_s + b_r$ block transfers plus $2 * b_r$ seeks
 - Each block in the inner relation s is read once for each block in the outer relation
- Best case ?
- In block nested-loop, use $M - 2$ disk blocks as blocking unit for outer relations, where M = memory size in blocks; use remaining two blocks to buffer inner relation and output
 - Cost = $\left\lceil \frac{b_r}{M-2} \right\rceil \times b_s + b_r$ block transfers + $2 \left\lceil \frac{b_r}{M-2} \right\rceil$ seeks
- If equi-join attribute forms a key or inner relation, stop inner loop on first match
- Scan inner loop forward and backward alternately, to make use of the blocks remaining in buffer (with LRU replacement)