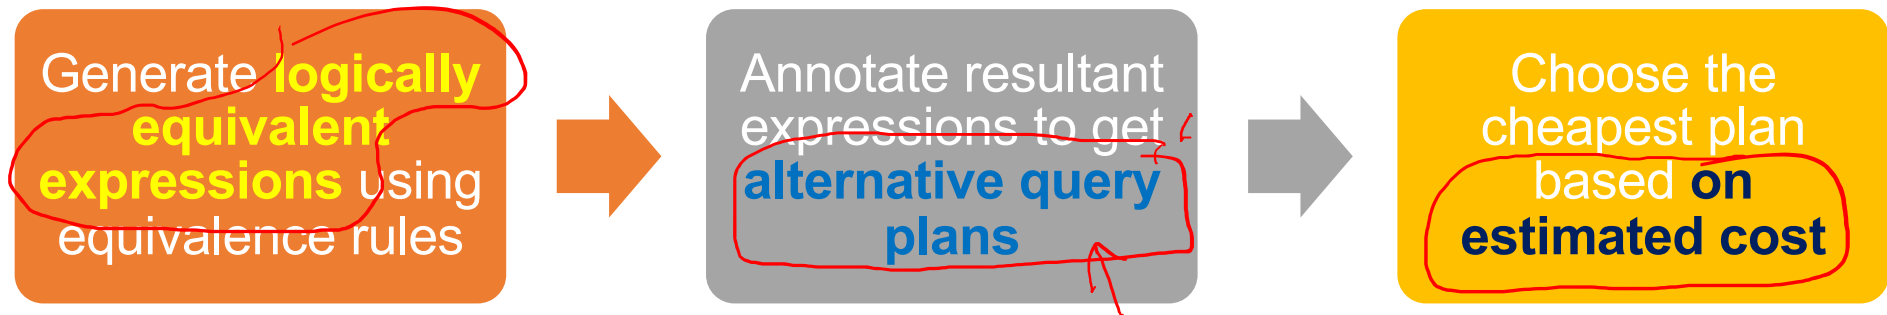


# **COL 362 & COL 632**

Query Optimization

31 Mar 2023

# Cost Based Query Optimization



- Estimation of plan cost based on
  - Statistical information about relations ✓
  - Statistics estimation for intermediate results ✓
  - Cost formulae for algorithms, computed using statistics ✓

# Transformation of Relational Expressions

- Two relational algebra expressions are said to be **equivalent** if the two expressions generate the same set of tuples on every *legal* database instance
- In SQL, inputs and outputs are multisets of tuples
  - Two expressions in the multiset version of the relational algebra are said to be equivalent if the two expressions generate the same multiset of tuples on every legal database instance.
- An **equivalence rule** says that expressions of two forms are equivalent
  - Can replace expression of first form by second, or vice versa

# Equivalence Rules

1. Conjunctive selection operations can be deconstructed into a sequence of individual selections.

$$\sigma_{\theta_1 \wedge \theta_2}(E) \equiv \sigma_{\theta_1}(\sigma_{\theta_2}(E))$$

2. Selection operations are commutative.

$$\sigma_{\theta_1}(\sigma_{\theta_2}(E)) \equiv \sigma_{\theta_2}(\sigma_{\theta_1}(E))$$

3. Only the last in a sequence of projection operations is needed, the others can be omitted.

$$\Pi_{L_1}(\Pi_{L_2}(\dots(\Pi_{L_n}(E))\dots)) \equiv \Pi_{L_1}(E)$$

where  $L_1 \subseteq L_2 \dots \subseteq L_n$

4. Selections can be combined with Cartesian products and theta joins.

a.  $\sigma_{\theta}(E_1 \times E_2) \equiv E_1 \bowtie_{\theta} E_2$

b.  $\sigma_{\theta_1}(E_1 \bowtie_{\theta_2} E_2) \equiv E_1 \bowtie_{\theta_1 \wedge \theta_2} E_2$

**Many more .... Refer to the textbook**

# Join Ordering

$$r_1 \bowtie r_2 \bowtie r_3$$

$$(r_1 \bowtie (r_2 \bowtie r_3))$$

- For all relations  $r_1, r_2$ , and  $r_3$ ,



$$(r_1 \bowtie r_2) \bowtie r_3 = r_1 \bowtie (r_2 \bowtie r_3)$$

(Join Associativity)

- If  $r_2 \bowtie r_3$  is quite large and  $r_1 \bowtie r_2$  is small, we choose

$$(r_1 \bowtie r_2) \bowtie r_3$$

so that we compute and store a smaller temporary relation.



# Join Ordering

- Consider the following two *equivalent expressions*

$$\Pi_{name, title} \left( \left( \sigma_{deptname="Music"} (instructor) \bowtie teaches \right) \bowtie \Pi_{courseid, title} (course) \right)$$

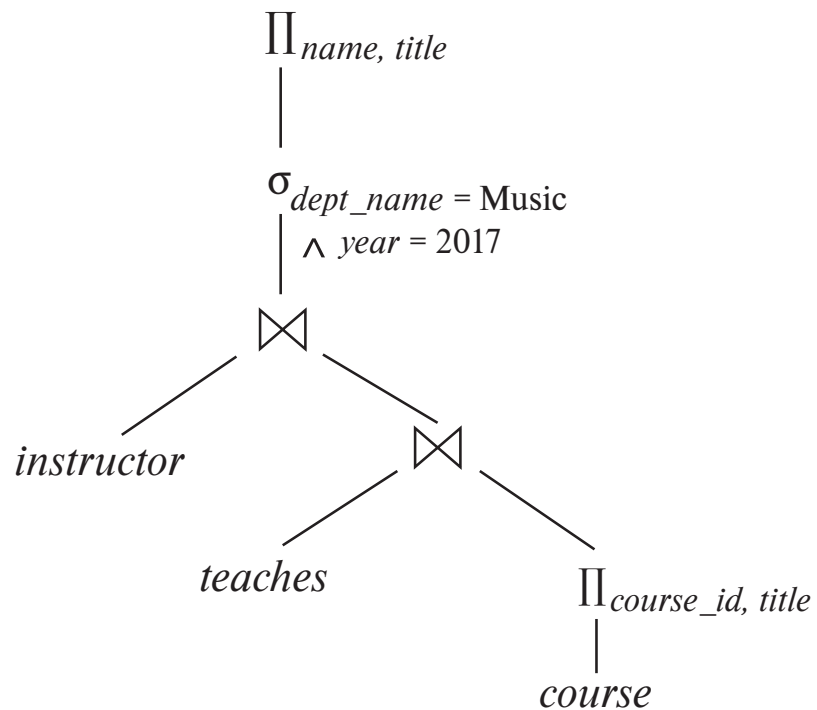
⑤

$$\Pi_{name, title} \left( \left( teaches \bowtie \Pi_{courseid, title} (course) \right) \bowtie \sigma_{deptname="Music"} (instructor) \right)$$

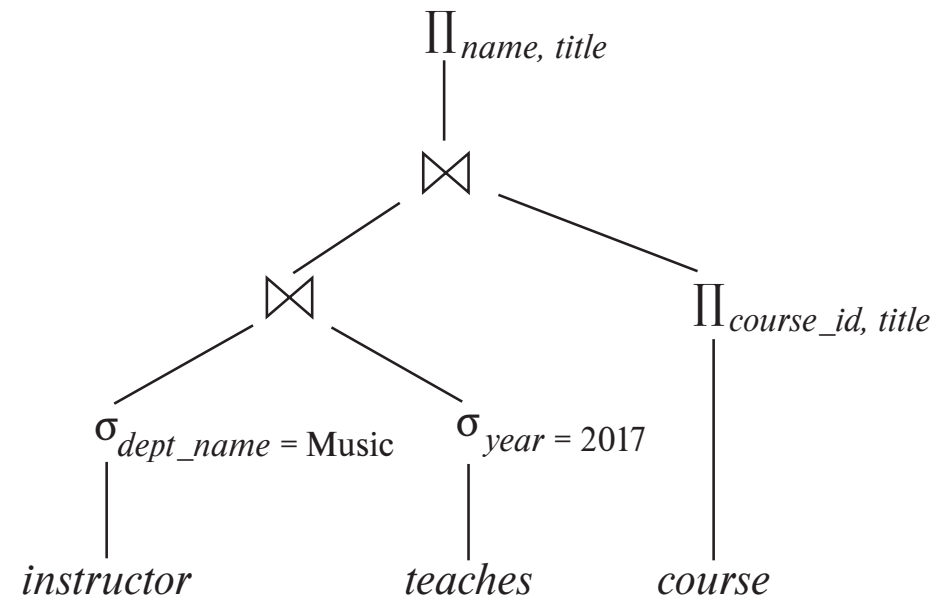
⑥

- Which is better?

# Join Ordering



(a) Initial expression tree



(b) Tree after multiple transformations

# Enumeration of Equivalent Expressions

- Query optimizers use equivalence rules to **systematically** generate expressions equivalent to the given expression

- Can generate all equivalent expressions

Repeat


- apply all applicable equivalence rules on every subexpression of every equivalent expression found so far
- add newly generated expressions to the set of equivalent expressions

Until no new equivalent expressions are generated above

- The above approach is very expensive in space and time
  1. Optimized plan generation based on transformation rules
  2. Special case approach for queries with only selections, projections and joins



# Cost-Based Optimization



- Consider finding the best join-order for  $r_1 \bowtie r_2 \bowtie \dots \bowtie r_n$
- The total number of possibilities is given by  $n-1$ -th Catalan number (# of binary trees with  $n$  leaves)

$$C(n) = \frac{1}{n+1} \binom{2n}{n} = \frac{(2n)!}{(n+1)!n!}$$

- No. of join orders of bushy trees (specific kind) :  $\frac{(2(n-1))!}{(n-1)!}$


$$\frac{(2n-2)!}{(n-1)!}$$

If  $n = 7 \Rightarrow 665,280$  join orders,

If  $n = 10 \Rightarrow 176$  billion join orders

# Transformation Based Optimization

- Space requirements reduced by sharing common sub-expressions:



- Time requirements are reduced by not generating all expressions

# Dynamic Programming in Optimization

- To find best join tree plan for a set  $S$  of  $n$  relations, consider all possible plans of the form:  $S_1$   $\bowtie$   $(S - S_1)$  where  $S_1$  is any non-empty subset of  $S$ .
- Recursively compute costs for joining subsets of  $S$  to find the cost of each plan. Choose the cheapest of the  $2^n - 2$  alternatives.
- Base case for recursion: single relation access plan
  - Apply all selections on  $R_i$  using best choice of indices on  $R_i$
- When plan for any subset is computed, store it and reuse it when it is required again, instead of recomputing it


# Join Order Optimization Algorithm

```
procedure findbestplan(S)
  if (bestplan[S].cost  $\neq \infty$ )
    return bestplan[S]
  // else bestplan[S] has not been computed earlier, compute it now
  if (S contains only 1 relation)
    set bestplan[S].plan and bestplan[S].cost based on the best way of accessing S
    /* Using selections on S and indices on S */
  else for each non-empty subset S1 of S such that S1  $\neq S$ 
    P1 = findbestplan(S1)
    P2 = findbestplan(S - S1)
    A = best algorithm for joining results of P1 and P2
    cost = P1.cost + P2.cost + cost of A
    if cost < bestplan[S].cost
      bestplan[S].cost = cost
      bestplan[S].plan = "execute P1.plan; execute P2.plan; join results of P1 and P2
        using A"
  return bestplan[S]
```


Sum is not always true

# Left Deep Join Trees

- In **left-deep join trees**, the right-hand-side input for each join is a relation, not the result of an intermediate join.




(a) Left-deep join tree



(b) Non-left-deep join tree

# Cost of Optimization

- With dynamic programming time complexity of optimization with bushy trees is  $O(3^n)$ .
  - With  $n = 10$ , this number is 59000 instead of 176 billion!
- Space complexity is  $O(2^n)$
- To find best left-deep join tree for a set of  $n$  relations:
  - Consider  $n$  alternatives with one relation as right-hand side input and the other relations as left-hand side input.
  - Replace “for each non-empty subset  $S_1$  of  $S$  such that  $S_1 \neq S$ ”
  - By: for each relation  $r$  in  $S$  let  $S_1 = S - r$ .
- If only left-deep trees are considered, time complexity of finding best join order is  $O(n 2^n)$  
  - Space complexity remains at  $O(2^n)$

# Interesting Sort Orders

- Consider the expression  $(r_1 \bowtie r_2) \bowtie r_3$  (with  $A$  as common attribute)
- An **interesting sort order** is a particular sort order of tuples that could be useful for a later operation
  - Using **merge-join** to compute  $r_1 \bowtie r_2$  may be costlier than **hash join** but the *result is sorted on  $A$*
  - Which in turn may make **merge-join with  $r_3$**  cheaper, which may reduce cost of join with  $r_3$  and minimizing overall cost
  - Sort order may also be useful for order by and for grouping
- Not sufficient to find the best join order for each subset of the set of  $n$  given relations
  - must find the best join order for each subset, **for each interesting sort order**
- Simple extension of earlier dynamic programming algorithms
  - Usually, number of interesting orders is small and doesn't affect time/space complexity significantly

# Heuristic Optimization

- Cost-based optimization is expensive, even with dynamic programming.
- Systems may use *heuristics* to reduce the number of choices that must be made in a cost-based fashion.
- Heuristic optimization transforms the query-tree by using a set of rules that typically (but not in all cases) improve execution performance:
  - Perform selection early (reduces the number of tuples)
  - Perform projection early (reduces the number of attributes)
  - Perform most restrictive selection and join operations (i.e. with smallest result size) before other similar operations.
  - Some systems use only heuristics, others combine heuristics with partial cost-based optimization.