

# COL362-632: Concurrency Control Schemes

---

Garima Gaur

April 11, 2023

- Disk block: DB is logically partitioned into blocks residing on disk.
- Buffer block: During a transaction execution, relevant blocks of DB are copied to the RAM.
- Each transaction has a local variable copy.

# Buffer Management

- Disk block: DB is logically partitioned into blocks residing on disk.
- Buffer block: During a transaction execution, relevant blocks of DB are copied to the RAM.
- Each transaction has a local variable copy.
- Transferring blocks between buffer and disk: input(B) and output(B).
- Read(b) includes input(B) and Write(b) copies local variable value to buffer block.

- Responsibility of DBMS to control the interaction of concurrent transactions.
  - Lock based Protocol
  - Timestamp based Protocol
  - Validation Based Protocol

# Locking Mechanism

- Locks to ensure sharing of data items is mutually exclusive.
- Two types of locks
  - Shared lock S(A): write operations not permitted
  - Exclusive lock X(A): both read and write operation permitted
- Lock compatibility matrix

	S	X
S	T	F
X	F	F

# Example

**T<sub>1</sub>**

---

```
lock-X(A);  
read(A);  
A := A - 50;  
write(A);  
unlock(A);  
lock-X(B);  
read(B);  
B := B + 50;  
write(B);  
unlock(B);
```

---

**T<sub>2</sub>**

---

```
lock-S(B);  
read(B);  
unlock(B);  
lock-S(A);  
read(A);  
unlock(A);  
display(A+B);
```

---

**T<sub>1</sub>**

---

```
lock-X(A)  
read(A)
```

```
write(A)  
unlock(A)  
lock-X(B)  
read(B)  
write(B)  
unlock(B)
```

**T<sub>2</sub>**

---

```
lock-S(B)  
read(B)  
unlock(B)
```

```
lock-S(A)  
read(A)  
unlock(A)
```

---

# Example

T <sub>1</sub>
lock-X(A);
read(A);
A := A - 50;
write(A);
unlock(A);
lock-X(B);
read(B);
B := B + 50;
write(B);
unlock(B);

T <sub>2</sub>
lock-S(B);
read(B);
unlock(B);
lock-S(A);
read(A);
unlock(A);
display(A+B);

T <sub>1</sub>	T <sub>2</sub>
lock-X(A)	
read(A)	
	lock-S(B)
	read(B)
	unlock(B)
write(A)	
unlock(A)	
lock-X(B)	
read(B)	
write(B)	
unlock(B)	
	lock-S(A)
	read(A)
	unlock(A)

- Only assessing the compatibility of locks is not a good measure.

# Does late unlocking help?

**T<sub>1</sub>**

---

lock-X(A);  
read(A);  
A := A - 50;  
write(A);  
lock-X(B);  
read(B);  
B := B + 50;  
write(B);  
unlock(A);  
unlock(B);

---

**T<sub>2</sub>**

---

lock-S(B);  
read(B);  
lock-S(A);  
read(A);  
display(A+B);  
unlock(B);  
unlock(A);

---

**T<sub>1</sub>**

---

lock-X(A)  
read(A)

write(A)  
lock-X(B)

**T<sub>2</sub>**

---

lock-S(B)  
read(B)

lock-S(A)

---



# Does late unlocking help?

**T<sub>1</sub>**

---

lock-X(A);  
read(A);  
A := A - 50;  
write(A);  
lock-X(B);  
read(B);  
B := B + 50;  
write(B);  
unlock(A);  
unlock(B);

---

**T<sub>2</sub>**

---

lock-S(B);  
read(B);  
lock-S(A);  
read(A);  
display(A+B);  
unlock(B);  
unlock(A);

---

**T<sub>1</sub>**

---

lock-X(A)  
read(A)  
  
write(A)  
lock-X(B)

---

**T<sub>2</sub>**

---

lock-S(B)  
read(B)  
  
lock-S(A)

---

- $T_1$  and  $T_2$  are in a **deadlock**!

## 2 Phase Locking Protocol

- Transactions operate in 2 phases – **growing phase** and **shrinking phase**.
- Growing phase – acquiring locks
- Shrinking phase – releasing locks
- Lock point – the point of acquiring last lock (end of growing phase).

## 2 Phase Locking Protocol

- Transactions operate in 2 phases – **growing phase** and **shrinking phase**.
- Growing phase – acquiring locks
- Shrinking phase – releasing locks
- Lock point – the point of acquiring last lock (end of growing phase).
- Locking point-based ordering of transactions is a serializability ordering of transactions.

- 2PL ensures conflict serializability but does not avoid deadlocks.
- Does not guarantee recoverability and cascadelessness.

# Variants of 2 PL

- 2PL ensures conflict serializability but does not avoid deadlocks.
- Does not guarantee recoverability and cascadelessness.
- Different variants of PL:
  - **Strict 2PL**: Avoids cascading rollbacks but may deadlock.
  - **Rigorous 2PL**: Another (more serial) variant to avoid cascading rollbacks.
  - **Conservative 2PL**: Deadlock free.

# Variants of 2 PL

- 2PL ensures conflict serializability but does not avoid deadlocks.
- Does not guarantee recoverability and cascadelessness.
- Different variants of PL:
  - **Strict 2PL**: Avoids cascading rollbacks but may deadlock.
  - **Rigorous 2PL**: Another (more serial) variant to avoid cascading rollbacks.
  - **Conservative 2PL**: Deadlock free.
- More concurrency comes with the possibility of deadlocks.
- Deadlock handling – prevention and detection-recovery.

# Timestamp based Protocol

- Deciding the ordering of transactions beforehand – assign timestamps to transactions.

# Timestamp based Protocol

- Deciding the ordering of transactions beforehand – assign timestamps to transactions.
- **TS**(**T<sub>i</sub>**): The timestamp assigned to a transaction  $T_i$  before it starts execution.
- $TS(T_i) < TS(T_j)$ :  $T_i$  has entered the system before the  $T_j$ .
- **W** – **TS**(**Q**): Largest timestamp of any transactions that executed write(Q).
- **R** – **TS**(**Q**): Largest timestamp of any transaction that executed read(Q).



# Timestamp-Ordering Protocol

- A transaction  $T_i$  issues *read*( $Q$ )
  - $\mathbf{TS}(T_i) < \mathbf{W} - \mathbf{TS}(Q)$ :  $T_i$  needs to read a value of  $Q$  that was already over-written. Rollback  $T_i$
  - $\mathbf{TS}(T_i) \geq \mathbf{W} - \mathbf{TS}(Q)$ : Read( $Q$ ) executed and set  $R - TS(Q) = \max\{R - TS(Q), TS(T_i)\}$

# Timestamp-Ordering Protocol

- A transaction  $T_i$  issues *read*( $Q$ )
  - $\mathbf{TS}(T_i) < \mathbf{W} - \mathbf{TS}(Q)$ :  $T_i$  needs to read a value of  $Q$  that was already over-written. Rollback  $T_i$
  - $\mathbf{TS}(T_i) \geq \mathbf{W} - \mathbf{TS}(Q)$ : Read( $Q$ ) executed and set  $R - TS(Q) = \max\{R - TS(Q), TS(T_i)\}$
- A transaction  $T_i$  issues *write*( $Q$ )
  - $\mathbf{TS}(T_i) < \mathbf{R} - \mathbf{TS}(Q)$ : The value of  $Q$  that  $T_i$  is producing was needed earlier, and the system assumed that the value would never be produced. Rollback  $T_i$ .
  - $\mathbf{TS}(T_i) < \mathbf{W} - \mathbf{TS}(Q)$ :  $T_i$  is trying to write an obsolete value of  $Q$ . Rollback  $T_i$

# Timestamp-Ordering Protocol

- Conflict serializable
- No deadlocks – processes are rolled back.

# Timestamp-Ordering Protocol

- Conflict serializable
- No deadlocks – processes are rolled back.

$T_4$	$T_5$
read(Q)	write(Q)
write(Q)	

- With timestamp ordering protocol, will  $T_4$  complete?

# Timestamp-Ordering Protocol

- Conflict serializable
- No deadlocks – processes are rolled back.

$T_4$	$T_5$
read(Q)	write(Q)
write(Q)	

- With timestamp ordering protocol, will  $T_4$  complete? **Obsolete write!**
- Write(Q) of  $T_4$  can be ignored. Why?

# Timestamp-Ordering Protocol

- Conflict serializable
- No deadlocks – processes are rolled back.

$T_4$	$T_5$
read(Q)	write(Q)
write(Q)	

- With timestamp ordering protocol, will  $T_4$  complete? **Obsolete write!**
- Write(Q) of  $T_4$  can be ignored. Why?
- Earlier transactions ( $TS(T_i) < TS(T_5)$ ) that attempt to read Q will be rolled back.
- Obsolete writes can be just ignored and transactions need not be roll-backed.

- Thomas' Write Rule: Modification in Timestamp ordering write operation protocol.
- $T_i$  issues write(Q),
  - $\mathbf{TS(T_i)} < \mathbf{R} - \mathbf{TS(Q)}$ : Rollback  $T_i$
  - $\mathbf{TS(T_i)} < \mathbf{W} - \mathbf{TS(Q)}$ :  $T_i$  is writing an obsolete value – **Ignore the write operation and proceed.**

## More concurrency

$T_4$	$T_5$
read(Q)	write(Q)
write(Q)	

- Is it conflict serializable? No.
- Is it view serializable?



## More concurrency

$T_4$	$T_5$
read(Q)	write(Q)
write(Q)	

- Is it conflict serializable? No.
- Is it view serializable? Equivalent to a serial schedule  $\langle T_4, T_5 \rangle$
- Thomas' write rule allows view serializable schedules.

- Majority of transactions are read-only transactions.
- Locking or ordering read-only is inefficient.
- Each transaction operates in 3 phases (in order)

- Majority of transactions are read-only transactions.
- Locking or ordering read-only is inefficient.
- Each transaction operates in 3 phases (in order)
  - **Read phase:** Read all data items and update only the local copies.
  - **Validation phase:** Use the validation test to determine if the transaction can proceed to the write phase without causing a violation of serializability.
  - **Write phase:** Updated local temporary values are copied to the database.

- Each transaction maintains 3 timestamp values:
  - **StartTS( $T_i$ )**: Time when  $T_i$  started execution.
  - **ValidateTS( $T_i$ )**: Time when  $T_i$  started its validation phase.
  - **FinishTS( $T_i$ )**: Time when  $T_i$  finished its Write phase.
- For serializability checks, *validateTS( $T_i$ )* of a transaction is treated as its timestamp *TS( $T_i$ )*.

# Validation Test

- Validation test for  $T_i$  requires all transactions  $T_k$  with  $TS(T_k) < TS(T_i)$  to satisfy one of the conditions,
  - $FinishTS(T_k) < StartTS(T_i)$ : Earlier transactions have already completed.
  - Set of data items written by  $T_k$  and set of items read by  $T_i$  does not overlap and  $FinishTS(T_k) < ValidationTS(T_i)$
- Validation-based protocols are called *optimistic* concurrency control.

# Snapshot Isolation

- [Snapshot Isolation](#): Oracle, PostgreSQL, SQL Server

# Snapshot Isolation

- **Snapshot Isolation**: Oracle, PostgreSQL, SQL Server
- Each transaction has its own snapshot of DB – reflecting changes made by the earlier committed transactions.
- Each transaction  $T_i$  has 2 timestamps:
  - $StartTS(T_i)$ : Start time of  $T_i$ .
  - $CommitTS(T_i)$ : Time at which validation is requested by  $T_i$ .

# Snapshot Isolation

- **Snapshot Isolation**: Oracle, PostgreSQL, SQL Server
- Each transaction has its own snapshot of DB – reflecting changes made by the earlier committed transactions.
- Each transaction  $T_i$  has 2 timestamps:
  - $StartTS(T_i)$ : Start time of  $T_i$ .
  - $CommitTS(T_i)$ : Time at which validation is requested by  $T_i$ .
- Each transaction  $T_i$  is unaware of updates of concurrent transaction  $T_j$ .
- Transactions are ordered:
  - *First committer wins*: If there exist a data item  $A$  that  $T_i$  intends to write and  $StartTS(T_i) < TS(A) < CommitTS(T_i)$ , then abort  $T_i$ .
  - *First updater wins*: Lock based approach, the transaction that acquires the item lock first, updates the item.



# Serializability Issue

$T_1$	$T_2$
read(A)	
read(B)	
	read(A)
	read(B)
A=B	
	B=A
write(A)	
	write(B)

# Serializability Issue

$T_1$	$T_2$
read(A)	
read(B)	
	read(A)
	read(B)
A=B	
	B=A
write(A)	
	write(B)

- Both transactions can commit under isolation snapshot – writing different data items.
- Is this schedule serializable?

# Serializability Issue

$T_1$	$T_2$
read(A)	
read(B)	
	read(A)
	read(B)
A=B	
	B=A
write(A)	
	write(B)

- Both transactions can commit under isolation snapshot – writing different data items.
- Is this schedule serializable? No; swapping values of A and B!

# Serializability Issue

$T_1$	$T_2$
read(A)	
read(B)	
	read(A)
	read(B)
A=B	
	B=A
write(A)	
	write(B)

- Both transactions can commit under isolation snapshot – writing different data items.
- Is this schedule serializable? No; swapping values of A and B!
- Serializable snapshot isolation (SSI): Details in Section 18.8.3 of the book.