

COL362-632: Recovery System

Garima Gaur

April 12, 2023

- Recovery system is responsible for atomicity and durability of transaction.

Recovery System

- Recovery system is responsible for atomicity and durability of transaction.
- Two tasks of a recovery system:
 - Record information about the transaction updates that will help in recovery after a crash.
 - After the system crashes, recover and ensure the durability of committed transactions.

Recovery System

- Recovery system is responsible for atomicity and durability of transaction.
- Two tasks of a recovery system:
 - Record information about the transaction updates that will help in recovery after a crash.
 - After the system crashes, recover and ensure the durability of committed transactions.
- **Log-based recovery**: Maintain a log that has a sequence of all the update operations of the transactions.
- Log is stored in stable store – information *never* gets lost.

- DB modification: Copying updated local variables to either the buffer blocks or the disk.

Database Modification

- DB modification: Copying updated local variables to either the buffer blocks or the disk.
- Two approaches of DB modification:
 - **Deferred modification**: Update DB only after the partial commit. (Validation protocol)
 - **Immediate modification**: Update DB along the instruction execution. (Lock-based protocol)

Database Modification

- DB modification: Copying updated local variables to either the buffer blocks or the disk.
- Two approaches of DB modification:
 - **Deferred modification**: Update DB only after the partial commit. (Validation protocol)
 - **Immediate modification**: Update DB along the instruction execution. (Lock-based protocol)
- Extra information collection prior DB modification by a recovery system.

Example

T₁	T₂
read(A);	read(C);
A := A - 50;	C = C - 100;
write(A);	write(C);
read(B);	
B := B + 50;	
write(B);	

Log	DB
$\langle T_1, Start \rangle$	
$\langle T_1, A, 1000, 950 \rangle$	
$\langle T_1, B, 2000, 2050 \rangle$	A = 950
	B = 2050
$\langle T_1, Commit \rangle$	
$\langle T_2, Start \rangle$	
$\langle T_2, C, 700, 600 \rangle$	
	C = 600
$\langle T_2, Commit \rangle$	

Example

		Log	DB
T₁		$\langle T_1, Start \rangle$	
read(A);		$\langle T_1, A, 1000, 950 \rangle$	
A := A - 50;	T₂	$\langle T_1, B, 2000, 2050 \rangle$	A = 950
write(A);	read(C);		B = 2050
read(B);	C = C - 100;	$\langle T_1, Commit \rangle$	
B := B + 50;	write(C);	$\langle T_2, Start \rangle$	
write(B);		$\langle T_2, C, 700, 600 \rangle$	
		$\langle T_2, Commit \rangle$	C = 600

- Each transaction T_i generates 2 types of log records:
 - For each write(Q) operation

$$\langle T_i, Q, Q_{old}, Q_{new} \rangle$$
 - Special records to denote the state of the transaction:

$$\langle T_i, Start \rangle; \langle T_i, Abort \rangle; \langle T_i, Commit \rangle$$

- A transaction commits after $\langle T_i, Commit \rangle$ is output to the log.

Redo and Undo Operation

- **Redo(T_i)**: Set the value of all the data items specified in the log record to the new updated value assigned by T_i .
- **Undo(T_i)**: Restore the value of all the data items updated by T_i to their old value.

Redo and Undo Operation

- **Redo(T_i)**: Set the value of all the data items specified in the log record to the new updated value assigned by T_i .
- **Undo(T_i)**: Restore the value of all the data items updated by T_i to their old value.

Log

$\langle T_1, \text{Start} \rangle$

$\langle T_1, A, 1000, 950 \rangle$

$\langle T_1, B, 2000, 2050 \rangle$

$\langle T_1, \text{Commit} \rangle$

$\langle T_2, \text{Start} \rangle$

$\langle T_2, C, 700, 600 \rangle$

$\langle T_2, \text{Commit} \rangle$

- Crash after write(A)

Redo and Undo Operation

- **Redo(T_i)**: Set the value of all the data items specified in the log record to the new updated value assigned by T_i .
- **Undo(T_i)**: Restore the value of all the data items updated by T_i to their old value.

Log

$\langle T_1, Start \rangle$

$\langle T_1, A, 1000, 950 \rangle$

$\langle T_1, B, 2000, 2050 \rangle$

$\langle T_1, Commit \rangle$

$\langle T_2, Start \rangle$

$\langle T_2, C, 700, 600 \rangle$

$\langle T_2, Commit \rangle$

- Crash after write(A) – undo(T_1).
- Crash after write(C) –

Redo and Undo Operation

- **Redo(T_i)**: Set the value of all the data items specified in the log record to the new updated value assigned by T_i .
- **Undo(T_i)**: Restore the value of all the data items updated by T_i to their old value.

Log

$\langle T_1, \text{Start} \rangle$

$\langle T_1, A, 1000, 950 \rangle$

$\langle T_1, B, 2000, 2050 \rangle$

$\langle T_1, \text{Commit} \rangle$

$\langle T_2, \text{Start} \rangle$

$\langle T_2, C, 700, 600 \rangle$

$\langle T_2, \text{Commit} \rangle$

- Crash after write(A) – undo(T_1).
- Crash after write(C) – redo(T_i) and undo(T_2).

Redo and Undo Operation

Log

$\langle T_1, \text{Start} \rangle$

$\langle T_1, A, 1000, 950 \rangle$

$\langle T_1, B, 2000, 2050 \rangle$

$\langle T_1, \text{Commit} \rangle$

$\langle T_2, \text{Start} \rangle$

$\langle T_2, C, 700, 600 \rangle$

$\langle T_2, \text{Commit} \rangle$

- Crash after write(B) operation – redo(T_1) or undo(T_1)?

Redo and Undo Operation

Log

$\langle T_1, \text{Start} \rangle$

$\langle T_1, A, 1000, 950 \rangle$

$\langle T_1, B, 2000, 2050 \rangle$

$\langle T_1, \text{Commit} \rangle$

$\langle T_2, \text{Start} \rangle$

$\langle T_2, C, 700, 600 \rangle$

$\langle T_2, \text{Commit} \rangle$

- Crash after write(B) operation – redo(T_1) or undo(T_1)?
- No commit record implies no persistence of change.

Checkpointing

- Recovery involves redoing and undoing transactions based on if $\langle T_{\bullet}, Commit \rangle$ or $\langle T_{\bullet}, Abort \rangle$ record is in the log.
- Logs can be enormous – longer search time and unnecessary Redo operations.

Checkpointing

- Recovery involves redoing and undoing transactions based on if $\langle T_{\bullet}, Commit \rangle$ or $\langle T_{\bullet}, Abort \rangle$ record is in the log.
- Logs can be enormous – longer search time and unnecessary Redo operations.
- Checkpointing is performed as:
 - Output all log records to stable storage.
 - Output all modified buffer blocks to the disk.
 - Make an entry $\langle checkpoint, L \rangle$ where L is the list of active transactions at checkpointing.
- No transaction performs any update actions during checkpointing.

Checkpointing

- Recovery involves redoing and undoing transactions based on if $\langle T_i, Commit \rangle$ or $\langle T_i, Abort \rangle$ record is in the log.
- Logs can be enormous – longer search time and unnecessary Redo operations.
- Checkpointing is performed as:
 - Output all log records to stable storage.
 - Output all modified buffer blocks to the disk.
 - Make an entry $\langle checkpoint, L \rangle$ where L is the list of active transactions at checkpointing.
- No transaction performs any update actions during checkpointing.
- All the transactions T_i with $\langle T_i, Commit \rangle$ after the checkpoint are redone.
- All the transactions without commit or abort record after the checkpoint are undone.

Recovery Algorithm

- Recovery takes place in 2 phases: redo phase and undo phase.
- **Redo Phase:** Scan the log forward from the last checkpoint
 - Initialize undo-list with L .
 - Whenever a non-special log record is encountered, the operation is redone.
 - If $\langle T_i, \text{Start} \rangle$ is found, add T_i to undo-list.
 - If either $\langle T_i, \text{Commit} \rangle$ or $\langle T_i, \text{Abort} \rangle$ is found, remove T_i from the undo-list.
- Undo-list is the list of all incomplete transactions.

Recovery Algorithm

- Recovery takes place in 2 phases: redo phase and undo phase.
- **Redo Phase:** Scan the log forward from the last checkpoint
 - Initialize undo-list with L .
 - Whenever a non-special log record is encountered, the operation is redone.
 - If $\langle T_i, \text{Start} \rangle$ is found, add T_i to undo-list.
 - If either $\langle T_i, \text{Commit} \rangle$ or $\langle T_i, \text{Abort} \rangle$ is found, remove T_i from the undo-list.
- Undo-list is the list of all incomplete transactions.
- **Undo Phase:** Scan the log backward from the end and focus on only undo-list transaction records.
 - When a non-special log record is found of undo list transaction T_i , perform undo operation as if T_i has failed and been rolled back.
 - If $\langle T_i, \text{Start} \rangle$ is found, writes $\langle T_i, \text{Abort} \rangle$ record and remove T_i from undo-list.
- Undo phase ends when undo list becomes empty.

Transaction Rollback

- T_i is rolled back during normal execution:
- Scan the log backward and for each record $\langle T_i, X, X_{old}, X_{new} \rangle$:
 - Update value of X to X_{old} .
 - Write a special redo-only record $\langle T_i, X, X_{old} \rangle$ to the log.
- When $\langle T_i, Start \rangle$ is found, stop scanning and write $\langle T_i, Abort \rangle$ to the log.

Example

$\langle T_0, \text{Start} \rangle$
 $\langle T_0, B, 2000, 2050 \rangle$
 $\langle T_1, \text{Start} \rangle$
 $\langle \text{checkpoint}, \{T_0, T_1\} \rangle$
 $\langle T_1, C, 700, 600 \rangle$
 $\langle T_1, \text{Commit} \rangle$
 $\langle T_2, \text{Start} \rangle$
 $\langle T_2, A, 500, 400 \rangle$
 $\langle T_0, B, 2000 \rangle$
 $\langle T_0, \text{Abort} \rangle$
crash
 \vdots
 $\langle T_2, A, 500 \rangle$
 $\langle T_2, \text{Abort} \rangle$