

# COL362-632: Transaction Management

---

April 5, 2023

# Transaction

- A logical unit of a program that reads and possibly updates data items.
- Data item at any granularity: token, tuple, relation

# Transaction

- A logical unit of a program that reads and possibly updates data items.
- Data item at any granularity: token, tuple, relation


$A = 1000$  and  $B = 950$

---

```
read(A);  
A := A - 50;  
write(A);  
read(B);  
B := B + 50;  
write(B);
```


---

# Life cycle of a transaction




- **Active**: Executing instructions.

# Life cycle of a transaction




- **Active**: Executing instructions.
- **Partially Committed**: All instructions executed, waiting for extra information and changes to be written on disk (DB).

# Life cycle of a transaction




- **Active**: Executing instructions.
- **Partially Committed**: All instructions executed, waiting for extra information and changes to be written on disk (DB).
- **Committed**: After the changes reflected in DB.

# Life cycle of a transaction



- **Active**: Executing instructions.
- **Partially Committed**: All instructions executed, waiting for extra information and changes to be written on disk (DB).
- **Committed**: After the changes reflected in DB.
- **Failed**: Transaction failed due to a failure/error.

# Life cycle of a transaction



- **Active**: Executing instructions.
- **Partially Committed**: All instructions executed, waiting for extra information and changes to be written on disk (DB).
- **Committed**: After the changes reflected in DB.
- **Failed**: Transaction failed due to a failure/error.
- **Aborted**: Transaction has failed and is rolled backed.



# Potential Issues

---

```
read(A);  
A := A - 50;  
write(A);  
read(B);  
B := B + 50;  
write(B);  
<commit>
```

---

- Transaction fails after debiting A – money lost!

# Potential Issues

---

```
read(A);  
A := A - 50;  
write(A);  
read(B);  
B := B + 50;  
write(B);  
<commit>
```

---

- Transaction fails after debiting A – money lost!
- Logical error:  $A := A + 50$  – application specified consistency violated!

# Potential Issues

---

```
read(A);  
A := A - 50;  
write(A);  
read(B);  
B := B + 50;  
write(B);  
<commit>
```

---

- Transaction fails after debiting A – money lost!
- Logical error:  $A := A + 50$  – application specified consistency violated!
- Another transaction is computing the average monthly balance of A – can generate the wrong value!

# Potential Issues

---

```
read(A);  
A := A - 50;  
write(A);  
read(B);  
B := B + 50;  
write(B);  
<commit>
```

---

- Transaction fails after debiting A – money lost!
- Logical error:  $A := A + 50$  – application specified consistency violated!
- Another transaction is computing the average monthly balance of A – can generate the wrong value!
- DB crashes after the transaction has commit operation: changes should not lost!

# Desirable characteristics of transaction

- **Atomicity**: Either all the operations in a transaction are executed or none at all.

# Desirable characteristics of transaction

- **Atomicity**: Either all the operations in a transaction are executed or none at all.
- **Consistency**: Execution of a transaction must preserve the consistency of the DB.

# Desirable characteristics of transaction

- **Atomicity**: Either all the operations in a transaction are executed or none at all.
- **Consistency**: Execution of a transaction must preserve the consistency of the DB.
- **Isolation**: Each transaction in the system is unaware of the other concurrent transactions.

# Desirable characteristics of transaction

- **Atomicity**: Either all the operations in a transaction are executed or none at all.
- **Consistency**: Execution of a transaction must preserve the consistency of the DB.
- **Isolation**: Each transaction in the system is unaware of the other concurrent transactions.
- **Durability**: The changes made by a successful transaction should persist against any system failure.

Popularly known as **ACID** properties



# Components responsible for ensuring these properties

- Atomicity: [Recovery System](#)
- Consistency: Responsibility of a programmer; Beyond data-integrity constraints;
- Isolation: [Concurrency control scheme](#)
- Durability: [Recovery System](#)

# Schedules

- **Schedule**: A chronological sequence of instructions of concurrent transactions.

# Schedules

- **Schedule**: A chronological sequence of instructions of concurrent transactions.

<b>T<sub>1</sub></b>	<b>T<sub>2</sub></b>
<hr/>	<hr/>
read(A);	read(A);
A := A - 50;	temp := A * 0.1;
write(A);	A := A - temp;
read(B);	write(A);
B := B + 50;	read(B);
write(B);	B := B + temp;
<hr/>	<hr/>

# Schedules

- **Schedule**: A chronological sequence of instructions of concurrent transactions.

<b>T<sub>1</sub></b>	<b>T<sub>2</sub></b>
read(A);	read(A);
A := A - 50;	temp := A * 0.1;
write(A);	A := A - temp;
read(B);	write(A);
B := B + 50;	read(B);
write(B);	B := B + temp;
	write(B);

	<b>T<sub>1</sub></b>	<b>T<sub>2</sub></b>
<i>i</i> <sub>0</sub>	read(A)	
<i>i</i> <sub>1</sub>		read(A)
<i>i</i> <sub>2</sub>		write(A)
<i>i</i> <sub>3</sub>	write(A)	
<i>i</i> <sub>4</sub>	read(B)	
<i>i</i> <sub>5</sub>	write(B)	
<i>i</i> <sub>6</sub>		read(B)
<i>i</i> <sub>7</sub>		write(B)

# Schedules

- **Schedule**: A chronological sequence of instructions of concurrent transactions.

<b>T<sub>1</sub></b>	<b>T<sub>2</sub></b>
read(A);	read(A);
A := A - 50;	temp := A * 0.1;
write(A);	A := A - temp;
read(B);	write(A);
B := B + 50;	read(B);
write(B);	B := B + temp;
	write(B);

	<b>T<sub>1</sub></b>	<b>T<sub>2</sub></b>
<i>i</i> <sub>0</sub>	read(A)	
<i>i</i> <sub>1</sub>		read(A)
<i>i</i> <sub>2</sub>		write(A)
<i>i</i> <sub>3</sub>	write(A)	
<i>i</i> <sub>4</sub>	read(B)	
<i>i</i> <sub>5</sub>	write(B)	
<i>i</i> <sub>6</sub>		read(B)
<i>i</i> <sub>7</sub>		write(B)

- Are concurrent transactions the same as concurrent processes?
  - ordering of Read and Write operations is crucial.

# Concurrent transactions

- $A = 1000$  and  $B = 950$

$T_1$	$T_2$
<u>read(A);</u> A := A - 50; write(A); read(B); B := B + 50; <u>write(B);</u>	<u>read(A);</u> temp := A * 0.1; A := A - temp; write(A); read(B); B := B + temp; <u>write(B);</u>

$T_1$	$T_2$
read(A)	read(A) write(A)
write(A) read(B) write(B)	read(B) write(B)

# Concurrent transactions

- $A = 1000$  and  $B = 950$

$T_1$	$T_2$
<code>read(A);</code> <code>A := A - 50;</code> <code>write(A);</code> <code>read(B);</code> <code>B := B + 50;</code> <code>write(B);</code>	<code>read(A);</code> <code>temp := A * 0.1;</code> <code>A := A - temp;</code> <code>write(A);</code> <code>read(B);</code> <code>B := B + temp;</code> <code>write(B);</code>

$T_1$	$T_2$
<code>read(A)</code>   <code>write(A)</code> <code>read(B)</code> <code>write(B)</code>	  <code>read(A)</code> <code>write(A)</code>   <code>read(B)</code> <code>write(B)</code>

- **Lost updates:** Changes made by a transaction are overwritten by another.

# Concurrent transactions

- $A = 1000$  and  $B = 950$

<u>T<sub>1</sub></u>	<u>T<sub>2</sub></u>
read(A); A := A - 50; write(A); read(B); B := B + 50; write(B);	read(A); temp := A * 0.1; A := A - temp; write(A); read(B); B := B + temp; write(B);

<u>T<sub>1</sub></u>	<u>T<sub>2</sub></u>
read(A)  write(A) read(B) write(B)	read(A) write(A)   read(B) write(B)

- **Lost updates:** Changes made by a transaction are overwritten by another.
- **Dirty reads:** A transaction reads the changes made by a failed transaction.



# Serial Schedules

- Serial schedules: Execute all instructions of a transaction followed by the instructions of the next transaction.

# Serial Schedules

- Serial schedules: Execute all instructions of a transaction followed by the instructions of the next transaction.
  - S1: T1 followed by T2
  - S2: T2 followed by T1

**T<sub>1</sub>**

---

```
read(A);  
A := A - 50;  
write(A);  
read(B);  
B := B + 50;  
write(B);
```

---

**T<sub>2</sub>**

---

```
read(A);  
temp := A * 0.1;  
A := A - temp;  
write(A);  
read(B);  
B := B + temp;  
write(B);
```

---

# Are serial schedules the real solution?

- Ensures transaction isolation perfectly.
- If execution is serial then the concurrency is not actually supported.

# Are serial schedules the real solution?

- Ensures transaction isolation perfectly.
- If execution is serial then the concurrency is not actually supported.
- Low throughput and under-utilized resources.
- Interleaving of instructions is desired.

## Requirement of “serial non-serial” schedules

# Requirement of “serial non-serial” schedules

<u>T<sub>1</sub></u>	<u>T<sub>2</sub></u>
read(A);	read(A);
A := A - 50;	temp := A * 0.1;
write(A);	A := A - temp;
read(B);	write(A);
B := B + 50;	read(B);
write(B);	B := B + temp;
	write(B);

<u>T<sub>1</sub></u>	<u>T<sub>2</sub></u>
read(A)	
write(A)	
	read(A)
	write(A)
read(B)	
write(B)	
	read(B)
	write(B)

## Requirement of “serial non-serial” schedules

<u>T<sub>1</sub></u>	<u>T<sub>2</sub></u>	<u>T<sub>1</sub></u>	<u>T<sub>2</sub></u>
read(A);	read(A);	read(A)	
A := A - 50;	temp := A * 0.1;	write(A)	
write(A);	A := A - temp;		read(A)
read(B);	write(A);		write(A)
B := B + 50;	read(B);	read(B)	
write(B);	B := B + temp;	write(B)	
	write(B);		read(B)
			write(B)

- Not all non-serial schedules lead to inconsistency!
- How to identify which schedules are “allowed”?

- A schedule is called *serializable* if it is **equivalent** to a serial schedule.



- A schedule is called *serializable* if it is **equivalent** to a serial schedule.
- Two types of equivalence:
  - **Conflict equivalence**
  - **View equivalence**

- A schedule is called *serializable* if it is **equivalent** to a serial schedule.
- Two types of equivalence:
  - **Conflict equivalence**
  - **View equivalence**
- A schedule S is conflict (view) serializable if it is conflict (view) equivalent to a serial schedule.

# Conflicting instructions

- Let I and J be two instructions from different transactions  $T_i$  and  $T_j$ :

# Conflicting instructions

- Let  $I$  and  $J$  be two instructions from different transactions  $T_i$  and  $T_j$ :
  1.  $I = R_i(X)$  and  $J = R_j(X)$ : No conflict.

# Conflicting instructions

- Let  $I$  and  $J$  be two instructions from different transactions  $T_i$  and  $T_j$ :
  1.  $I = R_i(X)$  and  $J = R_j(X)$ : No conflict.
  2.  $I = R_i(X)$  and  $J = W_j(X)$ : Value read by  $T_i$  is dependent on the order of  $I$  and  $J$ .

# Conflicting instructions

- Let  $I$  and  $J$  be two instructions from different transactions  $T_i$  and  $T_j$  :
  1.  $I = R_i(X)$  and  $J = R_j(X)$ : No conflict.
  2.  $I = R_i(X)$  and  $J = W_j(X)$ : Value read by  $T_i$  is dependent on the order of  $I$  and  $J$ .
  3.  $I = W_i(X)$  and  $J = R_j(X)$ : Similar as previous case.

# Conflicting instructions

- Let  $I$  and  $J$  be two instructions from different transactions  $T_i$  and  $T_j$  :
  1.  $I = R_i(X)$  and  $J = R_j(X)$ : No conflict.
  2.  $I = R_i(X)$  and  $J = W_j(X)$ : Value read by  $T_i$  is dependent on the order of  $I$  and  $J$ .
  3.  $I = W_i(X)$  and  $J = R_j(X)$ : Similar as previous case.
  4.  $I = W_i(X)$  and  $J = W_j(X)$ : Order of  $I$  and  $J$  decides the final value of  $X$

# Conflicting Instructions

- Two instructions I and J are said to be **conflicting** if
  - I and J are from two different transactions
  - Both operate on the same data item
  - At least one of them is a *write* operation



# Conflicting Instructions

- Two instructions I and J are said to be **conflicting** if
  - I and J are from two different transactions
  - Both operate on the same data item
  - At least one of them is a *write* operation

<b>T<sub>1</sub></b>	<b>T<sub>2</sub></b>
read(A)	
write(A)	
	read(A)
	write(A)
read(B)	
write(B)	
	read(B)
	write(B)

# Conflict Serializability

- Conflicting instructions impose a logical temporal ordering.
- Ordering of conflicting instructions should be preserved in an equivalent schedule.
- Consequently, non-conflicting instructions can be swapped.

# Conflict Serializability

- Conflicting instructions impose a logical temporal ordering.
- Ordering of conflicting instructions should be preserved in an equivalent schedule.
- Consequently, non-conflicting instructions can be swapped.
- If a schedule  $S$  can be transformed into a schedule  $S'$  by a series of swaps of non-conflicting instructions, then  $S$  and  $S'$  are **conflict equivalent**.

# Conflict Serializability

- Conflicting instructions impose a logical temporal ordering.
- Ordering of conflicting instructions should be preserved in an equivalent schedule.
- Consequently, non-conflicting instructions can be swapped.
- If a schedule  $S$  can be transformed into a schedule  $S'$  by a series of swaps of non-conflicting instructions, then  $S$  and  $S'$  are **conflict equivalent**.
- If  $S$  is *conflict* equivalent of a serial schedule, the  $S$  is said to be **conflict serializable**.

# View Serializability

- Two schedules are said to be view equivalent if they produce the same output (view).
- Two schedules  $S$  and  $S'$  are **view equivalent** if (for each data item  $X$ ),

# View Serializability

- Two schedules are said to be view equivalent if they produce the same output (view).
- Two schedules  $S$  and  $S'$  are **view equivalent** if (for each data item  $X$ ),
  1. Same transaction reads the initial value of  $X$  in both  $S$  and  $S'$ .

# View Serializability

- Two schedules are said to be view equivalent if they produce the same output (view).
- Two schedules  $S$  and  $S'$  are **view equivalent** if (for each data item  $X$ ),
  1. Same transaction reads the initial value of  $X$  in both  $S$  and  $S'$ .
  2. Same transaction performs the final write( $X$ ) operation in both schedules.

# View Serializability

- Two schedules are said to be view equivalent if they produce the same output (view).
- Two schedules  $S$  and  $S'$  are **view equivalent** if (for each data item  $X$ ),
  1. Same transaction reads the initial value of  $X$  in both  $S$  and  $S'$ .
  2. Same transaction performs the final write( $X$ ) operation in both schedules.
  3. If a transaction  $T_i$  reads a value of  $X$  produced by write( $X$ ) operation of transaction  $T_j$  in schedule  $S$ , then the same should hold in schedule  $S'$ .



## View and Conflict serializability relation

$T_1$	$T_2$	$T_3$
read(Q)	write(Q)	
write(Q)		write(Q)

- View equivalent to serial schedule  $S = \{T_1, T_2, T_3\}$

## View and Conflict serializability relation

$T_1$	$T_2$	$T_3$
read(Q)	write(Q)	
write(Q)		write(Q)

- View equivalent to serial schedule  $S = \{T_1, T_2, T_3\}$
- Is it conflict serializable?

## View and Conflict serializability relation

$T_1$	$T_2$	$T_3$
read(Q)	write(Q)	
write(Q)		write(Q)

- View equivalent to serial schedule  $S = \{T_1, T_2, T_3\}$
- Is it conflict serializable?
- Not conflict serializable because no pair of non-conflicting instructions!

Not all view serial schedules are conflict serializable, but all conflict serializable schedules are view serializable.

# Serializability Test

- Precedence graph: A directed graphs  $G(V, E)$  with  $V = \{T_i, \dots, T_j\}$  and and edge from transaction  $T_i$  to  $T_j$  exists if

# Serializability Test

- Precedence graph: A directed graphs  $G(V, E)$  with  $V = \{T_i, \dots, T_j\}$  and and edge from transaction  $T_i$  to  $T_j$  exists if
  - $R_i(X)$  precedes  $W_j(X)$

# Serializability Test

- Precedence graph: A directed graphs  $G(V, E)$  with  $V = \{T_i, \dots, T_j\}$  and and edge from transaction  $T_i$  to  $T_j$  exists if
  - $R_i(X)$  precedes  $W_j(X)$
  - $W_i(X)$  precedes  $R_j(X)$

# Serializability Test


- Precedence graph: A directed graphs  $G(V, E)$  with  $V = \{T_i, \dots, T_j\}$  and and edge from transaction  $T_i$  to  $T_j$  exists if
  - $R_i(X)$  precedes  $W_j(X)$
  - $W_i(X)$  precedes  $R_j(X)$
  - $W_i(X)$  precedes  $W_j(X)$

# Serializability Test

- Precedence graph: A directed graphs  $G(V, E)$  with  $V = \{T_i, \dots, T_j\}$  and an edge from transaction  $T_i$  to  $T_j$  exists if
  - $R_i(X)$  precedes  $W_j(X)$
  - $W_i(X)$  precedes  $R_j(X)$
  - $W_i(X)$  precedes  $W_j(X)$

$T_1$	$T_2$
read(A)	read(A) write(A)
write(A)	
read(B)	
write(B)	read(B) write(B)

- $read_1(A) < write_2(A)$
- $write_2(A) < write_1(A)$





- A schedule is conflict serializable *if and only if* the precedence graph is acyclic.

# Serializability Testing

- A schedule is conflict serializable *if and only if* the precedence graph is acyclic.
- Detecting a cycle in a graph:  $O(n+m)$  time.
- Topological sorting over the precedence graph provides the equivalent serial schedule.

# Serializability Testing

- A schedule is conflict serializable *if and only if* the precedence graph is acyclic.
- Detecting a cycle in a graph:  $O(n+m)$  time.
- Topological sorting over the precedence graph provides the equivalent serial schedule.
- Testing view serializability is NP-complete – no efficient algorithm exists.